



## Chapter 10. Analyzing the Meaning of Sentences

We have seen how useful it is to harness the power of a computer to process text on a large scale. However, now that we have the machinery of parsers and feature-based grammars, can we do anything similarly useful by analyzing the meaning of sentences? The goal of this chapter is to answer the following questions:

1. How can we represent natural language meaning so that a computer can process these representations?
2. How can we associate meaning representations with an unlimited set of sentences?
3. How can we use programs that connect the meaning representations of sentences to stores of knowledge?

Along the way we will learn some formal techniques in the field of logical semantics, and see how these can be used for interrogating databases that store facts about the world.

### Natural Language Understanding

#### QUERYING A DATABASE

Suppose we have a program that lets us type in a natural language question and gives us back the right answer:

*Example 10-1.*

1. Which country is Athens in?
2. Greece.

How hard is it to write such a program? And can we just use the same techniques



general way, we have to open up a whole new box of ideas and techniques, involving the representation of meaning.

So let's start off by assuming that we have data about cities and countries in a structured form. To be concrete, we will use a database table whose first few rows are shown in [Table 10-1](#).

NOTE

The data illustrated in [Table 10-1](#) is drawn from the Chat-80 system (Warren & Pereira, 1982). Population figures are given in thousands, but note that the data used in these examples dates back at least to the 1980s, and was already somewhat out of date at the point when (Warren & Pereira, 1982) was published.

Table 10-1. city\_table: A table of cities, countries, and populations

City	Country	Population
athens	greece	1368
bangkok	thailand	1178
barcelona	spain	1280
berlin	east_germany	3481
birmingham	united_kingdom	1112

The obvious way to retrieve answers from this tabular data involves writing queries in a database query language such as SQL.

NOTE

SQL (Structured Query Language) is a language designed for retrieving and managing data in relational databases. If you want to find out more about SQL, <http://www.w3schools.com/sql/> (<http://www.w3schools.com/sql/>) is a convenient online reference.

For example, executing the query [Example 10-2](#) will pull out the value 'greece':



This specifies a result set consisting of all values for the column `Country` in data rows where the value of the `City` column is `'athens'`.

How can we get the same effect using English as our input to the query system? The feature-based grammar formalism described in [Chapter 9](#) makes it easy to translate from English to SQL. The grammar *sql0.fcfg* illustrates how to assemble a meaning representation for a sentence in tandem with parsing the sentence. Each phrase structure rule is supplemented with a recipe for constructing a value for the feature `SEM`. You can see that these recipes are extremely simple; in each case, we use the string concatenation operation `+` to splice the values for the child constituents to make a value for the parent constituent.

```
>>> nltk.data.show_cfg('grammars/book_grammars/sql0.fcfg')
% start S
S[SEM=?np + WHERE + ?vp] -> NP[SEM=?np] VP[SEM=?vp]
VP[SEM=?v + ?pp] -> IV[SEM=?v] PP[SEM=?pp]
VP[SEM=?v + ?ap] -> IV[SEM=?v] AP[SEM=?ap]
NP[SEM=?det + ?n] -> Det[SEM=?det] N[SEM=?n]
PP[SEM=?p + ?np] -> P[SEM=?p] NP[SEM=?np]
AP[SEM=?pp] -> A[SEM=?a] PP[SEM=?pp]
NP[SEM='Country="greece"'] -> 'Greece'
NP[SEM='Country="china"'] -> 'China'
Det[SEM='SELECT'] -> 'Which' | 'What'
N[SEM='City FROM city_table'] -> 'cities'
IV[SEM=''] -> 'are'
A[SEM=''] -> 'located'
P[SEM=''] -> 'in'
```

This allows us to parse a query into SQL:

```
>>> from nltk import load_parser
>>> cp = load_parser('grammars/book_grammars/sql0.fcfg')
>>> query = 'What cities are located in China'
>>> trees = cp.nbest_parse(query.split())
>>> answer = trees[0].node['SEM']
>>> q = ' '.join(answer)
>>> print q
SELECT City FROM city_table WHERE Country="china"
```

NOTE

**Your Turn:** Run the parser with maximum tracing on, i.e., `cp = load_parser('grammars/book_grammars/sql0.fcfg', trace=3)`, and examine how the values of `SEM` are built up as complete edges are added to the chart.

Finally, we execute the query over the database `city.db` and retrieve some results:

```
>>> from nltk.sem import chat80
>>> rows = chat80.sql_query('corpora/city_database/city.db', q)
>>> for r in rows: print r[0], ❶
canton chungking dairen harbin kowloon mukden peking shanghai sia
```

Since each row `r` is a one-element tuple, we print out the member of the tuple rather than the tuple itself ❶.

To summarize, we have defined a task where the computer returns useful data in response to a natural language query, and we implemented this by translating a small subset of English into SQL. We can say that our NLTK code already “understands” SQL, given that Python is able to execute SQL queries against a database, and by extension it also “understands” queries such as *What cities are*



speaker of English, and have started to learn Dutch. Your teacher asks if you understand what [Example 10-3](#) means:

---

*Example 10-3.*

Margrietje houdt van Bruno.

If you know the meanings of the individual words in [Example 10-3](#), and know how these meanings are combined to make up the meaning of the whole sentence, you might say that [Example 10-3](#) means the same as *Margrietje loves Bruno*.

An observer—let’s call her Olga—might well take this as evidence that you do grasp the meaning of [Example 10-3](#). But this would depend on Olga herself understanding English. If she doesn’t, then your translation from Dutch to English is not going to convince her of your ability to understand Dutch. We will return to this issue shortly.

The grammar *sql0.fcfg*, together with the NLTK Earley parser, is instrumental in carrying out the translation from English to SQL. How adequate is this grammar? You saw that the SQL translation for the whole sentence was built up from the translations of the components. However, there does not seem to be a lot of justification for these component meaning representations. For example, if we look at the analysis of the noun phrase *Which cities*, the determiner and noun correspond respectively to the SQL fragments `SELECT` and `City FROM city_table`. But neither of these has a well-defined meaning in isolation from the other.

There is another criticism we can level at the grammar: we have “hard-wired” an embarrassing amount of detail about the database into it. We need to know the name of the relevant table (e.g., `city_table`) and the names of the fields. But our database could have contained exactly the same rows of data yet used a different table name and different field names, in which case the SQL queries would not be executable. Equally, we could have stored our data in a different format, such as XML, in which case retrieving the same results would require us to translate our English queries into an XML query language rather than SQL. These considerations suggest that we should be translating English into something that is more abstract and generic than SQL.

In order to sharpen the point, let’s consider another English query and its translation:

---

*Example 10-4.*

1. What cities are in China and have populations above 1,000,000?
2. `SELECT City FROM city_table WHERE Country = 'china' AND Population > 1000`

NOTE

**Your Turn:** Extend the grammar *sql0.fcfg* so that it will translate **a** into **b**, and check the values returned by the query. Remember that figures in the Chat-80 database are given in thousands, hence 1000 in (4b) represents one million inhabitants.

You will probably find it easiest to first extend the grammar to handle queries like *What cities have populations above 1,000,000* before tackling conjunction. After you have had a go at this task, you can compare your solution to `grammars/book_grammars/sql1.fcfg` in the NLTK data distribution.



Observe that the *and* conjunction in **a** is translated into an AND in the SQL

counterpart, **b**. The latter tells us to select results from rows where two

conditions are true together: the value of the Country column is 'china' and the value of the Population column is greater than 1000. This interpretation for *and* involves a new idea: it talks about *what is true in some particular situation*, and tells us that *Cond1 AND Cond2* is true in situation *s* if and only if condition *Cond1* is true in *s* and condition *Cond2* is true in *s*. Although this doesn't account for the full range of meanings of *and* in English, it has the nice property that it is independent of any query language. In fact, we have given it the standard interpretation from classical logic. In the following sections, we will explore an approach in which sentences of natural language are translated into logic instead of an executable query language such as SQL. One advantage of logical formalisms is that they are more abstract and therefore more generic. If we wanted to, once we had our translation into logic, we could then translate it into various other special-purpose languages. In fact, most serious attempts to query databases via natural language have used this methodology.

## NATURAL LANGUAGE, SEMANTICS, AND LOGIC

We started out trying to capture the meaning of **a** by translating it into a query in another language, SQL, which the computer could interpret and execute. But this still begged the question whether the translation was correct. Stepping back from database query, we noted that the meaning of *and* seems to depend on being able to specify when statements are true or not in a particular situation. Instead of translating a sentence *S* from one language to another, we try to say what *S* is *about* by relating it to a situation in the world. Let's pursue this further. Imagine there is a situation *s* where there are two entities, Margrietje and her favorite doll, Brunoke. In addition, there is a relation holding between the two entities, which we will call the *love* relation. If you understand the meaning of **Example 10-3**, then you know that it is true in situation *s*. In part, you know this because you know that *Margrietje* refers to Margrietje, *Brunoke* refers to Brunoke, and *houdt van* refers to the *love* relation.

We have introduced two fundamental notions in semantics. The first is that declarative sentences are *true or false in certain situations*. The second is that definite noun phrases and proper nouns *refer to things in the world*. So **Example 10-3** is true in a situation where Margrietje loves the doll Brunoke, here illustrated in **Figure 10-1**.



Figure 10-1. Depiction of a situation in which Margrietje loves Brunoke.

Once we have adopted the notion of truth in a situation, we have a powerful tool



7 cannot be. In other words, the sentences in [Example 10-5](#) are **consistent**, whereas those in [Example 10-6](#) and [Example 10-7](#) are **inconsistent**.

---

*Example 10-5.*

1. Sylvania is to the north of Freedonia.
2. Freedonia is a republic.

---

*Example 10-6.*

1. The capital of Freedonia has a population of 9,000.
2. No city in Freedonia has a population of 9,000.

---

*Example 10-7.*

1. Sylvania is to the north of Freedonia.
2. Freedonia is to the north of Sylvania.

We have chosen sentences about fictional countries (featured in the Marx Brothers' 1933 movie *Duck Soup*) to emphasize that your ability to reason about these examples does not depend on what is true or false in the actual world. If you know the meaning of the word *no*, and also know that the capital of a country is a city in that country, then you should be able to conclude that the two sentences in [Example 10-6](#) are inconsistent, regardless of where Freedonia is or what the population of its capital is. That is, there's no possible situation in which both sentences could be true. Similarly, if you know that the relation expressed by *to the north of* is asymmetric, then you should be able to conclude that the two sentences in [Example 10-7](#) are inconsistent.

Broadly speaking, logic-based approaches to natural language semantics focus on those aspects of natural language that guide our judgments of consistency and inconsistency. The syntax of a logical language is designed to make these features formally explicit. As a result, determining properties like consistency can often be reduced to symbolic manipulation, that is, to a task that can be carried out by a computer. In order to pursue this approach, we first want to develop a technique for representing a possible situation. We do this in terms of something that logicians call a “model.”

A **model** for a set  $W$  of sentences is a formal representation of a situation in which all the sentences in  $W$  are true. The usual way of representing models involves set theory. The domain  $D$  of discourse (all the entities we currently care about) is a set of individuals, while relations are treated as sets built up from  $D$ . Let's look at a concrete example. Our domain  $D$  will consist of three children, Stefan, Klaus, and Evi, represented respectively as  $s$ ,  $k$ , and  $e$ . We write this as  $D = \{s, k, e\}$ . The expression *boy* denotes the set consisting of Stefan and Klaus, the expression *girl* denotes the set consisting of Evi, and the expression *is running* denotes the set consisting of Stefan and Evi. [Figure 10-2](#) is a graphical rendering of the model.






Figure 10-2. Diagram of a model containing a domain  $D$  and subsets of  $D$  corresponding to the predicates *boy*, *girl*, and *is running*.

Later in this chapter we will use models to help evaluate the truth or falsity of English sentences, and in this way to illustrate some methods for representing meaning. However, before going into more detail, let's put the discussion into a broader perspective, and link back to a topic that we briefly raised in [Automatic Natural Language Understanding](#). Can a computer understand the meaning of a sentence? And how could we tell if it did? This is similar to asking "Can a computer think?" Alan Turing famously proposed to answer this by examining the ability of a computer to hold sensible conversations with a human (Turing, 1950). Suppose you are having a chat session with a person and a computer, but you are not told at the outset which is which. If you cannot identify which of your partners is the computer after chatting with each of them, then the computer has successfully imitated a human. If a computer succeeds in passing itself off as human in this "imitation game" (or "Turing Test" as it is popularly known), then according to Turing, we should be prepared to say that the computer *can* think and can be said to be intelligent. So Turing side-stepped the question of somehow examining the internal states of a computer by instead using its *behavior* as evidence of intelligence. By the same reasoning, we have assumed that in order to say that a computer understands English, it just needs to behave as though it did. What is important here is not so much the specifics of Turing's imitation game, but rather the proposal to judge a capacity for natural language understanding in terms of observable behavior.

### Propositional Logic

A logical language is designed to make reasoning formally explicit. As a result, it can capture aspects of natural language which determine whether a set of sentences is consistent. As part of this approach, we need to develop logical representations of a sentence  $\phi$  that formally capture the **truth-conditions** of  $\phi$ . We'll start off with a simple example:

*Example 10-8.*

[Klaus chased Evi] and [Evi ran away].

Let's replace the two sub-sentences in [Example 10-8](#) by  $\phi$  and  $\psi$  respectively, and put  $\&$  for the logical operator corresponding to the English word *and*:  $\phi \& \psi$ . This structure is the **logical form** of [Example 10-8](#).

**Propositional logic** allows us to represent just those parts of linguistic structure that correspond to certain sentential connectives. We have just looked at *and*.



**propositional symbols**, often written as  $P$ ,  $Q$ ,  $R$ , etc. There are varying conventions for representing Boolean operators. Since we will be focusing on ways of exploring logic within NLTK, we will stick to the following ASCII versions of the operators:

```
>>> nltk.boolean_ops()
negation      ~
conjunction    &
disjunction    |
implication    ->
equivalence    <->
```

From the propositional symbols and the Boolean operators we can build an infinite set of **well-formed formulas** (or just formulas, for short) of propositional logic. First, every propositional letter is a formula. Then if  $\phi$  is a formula, so is  $\neg\phi$ . And if  $\phi$  and  $\psi$  are formulas, then so are  $(\phi \ \& \ \psi)$ ,  $(\phi \ | \ \psi)$ ,  $(\phi \ -> \ \psi)$ , and  $(\phi \ <-> \ \psi)$ .

**Table 10-2** specifies the truth-conditions for formulas containing these operators. As before we use  $\phi$  and  $\psi$  as variables over sentences, and abbreviate *if and only if* as *iff*.

Table 10-2. Truth conditions for the Boolean operators in propositional logic

Boolean operator	Truth conditions		
negation ( <i>it is not the case that ...</i> )	$\neg\phi$ is true in $s$	iff	$\phi$ is false in $s$
conjunction ( <i>and</i> )	$(\phi \ \& \ \psi)$ is true in $s$	iff	$\phi$ is true in $s$ and $\psi$ is true in $s$
disjunction ( <i>or</i> )	$(\phi \   \ \psi)$ is true in $s$	iff	$\phi$ is true in $s$ or $\psi$ is true in $s$
implication ( <i>if ..., then ...</i> )	$(\phi \ -> \ \psi)$ is true in $s$	iff	$\phi$ is false in $s$ or $\psi$ is true in $s$
equivalence ( <i>if and only if</i> )	$(\phi \ <-> \ \psi)$ is true in $s$	iff	$\phi$ and $\psi$ are both true in $s$ or both false in $s$

These rules are generally straightforward, though the truth conditions for implication depart in many cases from our usual intuitions about the conditional in English. A formula of the form  $(P \ -> \ Q)$  is false only when  $P$  is true and  $Q$  is false. If  $P$  is false (say,  $P$  corresponds to *The moon is made of green cheese*) and  $Q$  is true (say,  $Q$  corresponds to *Two plus two equals four*), then  $P \ -> \ Q$  will come out true.

NLTK's `LogicParser()` parses logical expressions into various subclasses of





```
>>> lp = nltk.LogicParser()
>>> lp.parse('(P & Q)')
<NegatedExpression ~(P & Q)>
>>> lp.parse('P & Q')
<AndExpression (P & Q)>
>>> lp.parse('P | (R -> Q)')
<OrExpression (P | (R -> Q))>
>>> lp.parse('P <-> -- P')
<IffExpression (P <-> --P)>
```

From a computational perspective, logics give us an important tool for performing inference. Suppose you state that Freedonia is not to the north of Sylvania, and you give as your reasons that Sylvania is to the north of Freedonia. In this case, you have produced an **argument**. The sentence *Sylvania is to the north of Freedonia* is the **assumption** of the argument, while *Freedonia is not to the north of Sylvania* is the **conclusion**. The step of moving from one or more assumptions to a conclusion is called **inference**. Informally, it is common to write arguments in a format where the conclusion is preceded by *therefore*.

---

*Example 10-9.*

Sylvania is to the north of Freedonia.

Therefore, Freedonia is not to the north of Sylvania.

An argument is **valid** if there is no possible situation in which its premises are all true and its conclusion is not true.

Now, the validity of [Example 10-9](#) crucially depends on the meaning of the phrase *to the north of*, in particular, the fact that it is an asymmetric relation:

---

*Example 10-10.*

if  $x$  is to the north of  $y$  then  $y$  is not to the north of  $x$ .

Unfortunately, we can't express such rules in propositional logic: the smallest elements we have to play with are atomic propositions, and we cannot "look inside" these to talk about relations between individuals  $x$  and  $y$ . The best we can do in this case is capture a particular case of the asymmetry. Let's use the propositional symbol  $\text{SnF}$  to stand for *Sylvania is to the north of Freedonia* and  $\text{FnS}$  for *Freedonia is to the north of Sylvania*. To say that *Freedonia is not to the north of Sylvania*, we write  $\neg\text{FnS}$ . That is, we treat *not* as equivalent to the phrase *it is not the case that ...*, and translate this as the one-place Boolean operator  $\neg$ . Replacing  $x$  and  $y$  in [Example 10-10](#) by *Sylvania* and *Freedonia* respectively gives us an implication that can be written as:

---

*Example 10-11.*

$\text{SnF} \rightarrow \neg\text{FnS}$

How about giving a version of the complete argument? We will replace the first sentence of [Example 10-9](#) by two formulas of propositional logic:  $\text{SnF}$ , and also the implication in [Example 10-11](#), which expresses (rather poorly) our background knowledge of the meaning of *to the north of*. We'll write  $[A_1, \dots, A_n] / C$  to represent the argument that conclusion  $C$  follows from assumptions  $[A_1, \dots, A_n]$ . This leads to the following as a representation of argument [Example 10-9](#):

---

*Example 10-12.*

$[\text{SnF}, \text{SnF} \rightarrow \neg\text{FnS}] / \neg\text{FnS}$

This is a valid argument: if  $\text{SnF}$  and  $\text{SnF} \rightarrow \neg\text{FnS}$  are both true in a situation  $s$ , then  $\neg\text{FnS}$  must also be true in  $s$ . By contrast, if  $\text{FnS}$  were true, this would conflict with our understanding that two objects cannot both be to the north of each other in any possible situation. Equivalently, the list  $[\text{SnF}, \text{SnF} \rightarrow \neg$



Arguments can be tested for “syntactic validity” by using a proof system. We will say a little bit more about this later on in [First-Order Logic](#). Logical proofs can be carried out with NLTK’s `inference` module, for example, *via* an interface to the third-party theorem prover Prover9. The inputs to the inference mechanism first have to be parsed into logical expressions by `LogicParser()`.

```
>>> lp = nltk.LogicParser()
>>> SnF = lp.parse('SnF')
>>> NotFnS = lp.parse('-FnS')
>>> R = lp.parse('SnF -> -FnS')
>>> prover = nltk.Prover9()
>>> prover.prove(NotFnS, [SnF, R])
True
```

Here’s another way of seeing why the conclusion follows.  $SnF \rightarrow -FnS$  is semantically equivalent to  $\neg SnF \vee \neg FnS$ , where  $\vee$  is the two-place operator corresponding to *or*. In general,  $\varphi \vee \psi$  is true in a situation  $s$  if either  $\varphi$  is true in  $s$  or  $\psi$  is true in  $s$ . Now, suppose both  $SnF$  and  $\neg SnF \vee \neg FnS$  are true in situation  $s$ . If  $SnF$  is true, then  $\neg FnS$  cannot also be true; a fundamental assumption of classical logic is that a sentence cannot be both true and false in a situation. Consequently,  $\neg FnS$  must be true.

Recall that we interpret sentences of a logical language relative to a model, which is a very simplified version of the world. A model for propositional logic needs to assign the values `True` or `False` to every possible formula. We do this inductively: first, every propositional symbol is assigned a value, and then we compute the value of complex formulas by consulting the meanings of the Boolean operators (i.e., [Table 10-2](#)) and applying them to the values of the formula’s components. A `Valuation` is a mapping from basic symbols of the logic to their values. Here’s an example:

```
>>> val = nltk.Valuation([( 'P', True), ( 'Q', True), ( 'R', False)])
```

We initialize a `Valuation` with a list of pairs, each of which consists of a semantic symbol and a semantic value. The resulting object is essentially just a dictionary that maps logical symbols (treated as strings) to appropriate values.

```
>>> val['P']
True
```

As we will see later, our models need to be somewhat more complicated in order to handle the more complex logical forms discussed in the next section; for the time being, just ignore the `dom` and `g` parameters in the following declarations.

```
>>> dom = set([])
>>> g = nltk.Assignment(dom)
```

Now let’s initialize a model `m` that uses `val`:

```
>>> m = nltk.Model(dom, val)
```

Every model comes with an `evaluate()` method, which will determine the semantic value of logical expressions, such as formulas of propositional logic; of course, these values depend on the initial truth values we assigned to propositional symbols such as `P`, `Q`, and `R`.

```
>>> print m.evaluate('(P & Q)', g)
True
>>> print m.evaluate('(¬(P & Q))', g)
False
>>> print m.evaluate('(P & R)', g)
False
```



## NOTE

**Your Turn:** Experiment with evaluating different formulas of propositional logic. Does the model give the values that you expected?

Up until now, we have been translating our English sentences into propositional logic. Because we are confined to representing atomic sentences with letters such as  $P$  and  $Q$ , we cannot dig into their internal structure. In effect, we are saying that there is no semantic benefit in dividing atomic sentences into subjects, objects, and predicates. However, this seems wrong: if we want to formalize arguments such as [Example 10-9](#), we have to be able to “look inside” basic sentences. As a result, we will move beyond propositional logic to something more expressive, namely first-order logic. This is what we turn to in the next section.

## First-Order Logic

In the remainder of this chapter, we will represent the meaning of natural language expressions by translating them into first-order logic. Not all of natural language semantics can be expressed in first-order logic. But it is a good choice for computational semantics because it is expressive enough to represent many aspects of semantics, and on the other hand, there are excellent systems available off the shelf for carrying out automated inference in first-order logic.

Our next step will be to describe how formulas of first-order logic are constructed, and then how such formulas can be evaluated in a model.

## SYNTAX

First-order logic keeps all the Boolean operators of propositional logic, but it adds some important new mechanisms. To start with, propositions are analyzed into predicates and arguments, which takes us a step closer to the structure of natural languages. The standard construction rules for first-order logic recognize **terms** such as individual variables and individual constants, and **predicates** that take differing numbers of **arguments**. For example, *Angus walks* might be formalized as *walk(angus)* and *Angus sees Bertie* as *see(angus, bertie)*. We will call *walk* a **unary predicate**, and *see* a **binary predicate**. The symbols used as predicates do not have intrinsic meaning, although it is hard to remember this. Returning to one of our earlier examples, there is no *logical* difference between

$a$  and  $b$ .

*Example 10-13.*

1. *love(margrietje, brunoke)*
2. *houden\_van(margrietje, brunoke)*

By itself, first-order logic has nothing substantive to say about lexical semantics—the meaning of individual words—although some theories of lexical semantics can be encoded in first-order logic. Whether an atomic predication like *see(angus, bertie)* is true or false in a situation is not a matter of logic, but depends on the particular valuation that we have chosen for the constants *see*, *angus*, and *bertie*. For this reason, such expressions are called **non-logical constants**. By contrast, **logical constants** (such as the Boolean operators) always receive the same interpretation in every model for first-order logic.

We should mention here that one binary predicate has special status, namely equality, as in formulas such as *angus = aj*. Equality is regarded as a logical constant, since for individual terms  $t_1$  and  $t_2$ , the formula  $t_1 = t_2$  is true if and



It is often helpful to inspect the syntactic structure of expressions of first-order logic, and the usual way of doing this is to assign **types** to expressions. Following the tradition of Montague grammar, we will use two **basic types**:  $e$  is the type of entities, while  $t$  is the type of formulas, i.e., expressions that have truth values. Given these two basic types, we can form **complex types** for function expressions. That is, given any types  $\sigma$  and  $\tau$ ,  $\langle\sigma, \tau\rangle$  is a complex type corresponding to functions from ' $\sigma$  things' to ' $\tau$  things'. For example,  $\langle e, t\rangle$  is the type of expressions from entities to truth values, namely unary predicates. The `LogicParser` can be invoked so that it carries out type checking.

```
>>> tlp = nltk.LogicParser(type_check=True)
>>> parsed = tlp.parse('walk(angus)')
>>> parsed.argument
<ConstantExpression angus>
>>> parsed.argument.type
e
>>> parsed.function
<ConstantExpression walk>
>>> parsed.function.type
<e, ?>
```

Why do we see `<e, ?>` at the end of this example? Although the type-checker will try to infer as many types as possible, in this case it has not managed to fully specify the type of `walk`, since its result type is unknown. Although we are intending `walk` to receive type `<e, t>`, as far as the type-checker knows, in this context it could be of some other type, such as `<e, e>` or `<e, t>>`. To help the type-checker, we need to specify a **signature**, implemented as a dictionary that explicitly associates types with non-logical constants:

```
>>> sig = {'walk': '<e, t>'}
>>> parsed = tlp.parse('walk(angus)', sig)
>>> parsed.function.type
<e, t>
```

A binary predicate has type  $\langle e, \langle e, t\rangle\rangle$ . Although this is the type of something which combines first with an argument of type  $e$  to make a unary predicate, we represent binary predicates as combining directly with their two arguments. For example, the predicate *see* in the translation of *Angus sees Cyril* will combine with its arguments to give the result *see(angus, cyril)*.

In first-order logic, arguments of predicates can also be individual variables such as  $x$ ,  $y$ , and  $z$ . In NLTK, we adopt the convention that variables of type  $e$  are all lowercase. Individual variables are similar to personal pronouns like *he*, *she*, and *it*, in that we need to know about the context of use in order to figure out their denotation. One way of interpreting the pronoun in [Example 10-14](#) is by pointing to a relevant individual in the local context.

*Example 10-14.*

He disappeared.

Another way is to supply a textual antecedent for the pronoun *he*, for example, by uttering **a** prior to [Example 10-14](#). Here, we say that *he* is **coreferential** with the noun phrase *Cyril*. In such a context, [Example 10-14](#) is semantically equivalent to **b**.

*Example 10-15.*

1. Cyril is Angus's dog.
2. Cyril disappeared.

Consider by contrast the occurrence of *he* in **a**. In this case, it is **bound** by the indefinite NP *a dog*, and this is a different relationship than coreference. If we



replace the pronoun *he* by *a dog*, the result **b** is *not* semantically equivalent to

**a** .

*Example 10-16.*

1. Angus had a dog but he disappeared.
2. Angus had a dog but a dog disappeared.

Corresponding to **a** , we can construct an **open formula** **b** with two occurrences of the variable *x*. (We ignore tense to simplify exposition.)

*Example 10-17.*

1. He is a dog and he disappeared.
2. *dog(x) & disappear(x)*

By placing an **existential quantifier**  $\exists x$  (“for some *x*”) in front of **b** , we can **bind** these variables, as in **a** , which means **b** or, more idiomatically, **c** .

*Example 10-18.*

1.  $\exists x.(dog(x) \& disappear(x))$
2. At least one entity is a dog and disappeared.
3. A dog disappeared.

Here is the NLTK counterpart of **a** :

*Example 10-19.*

exists x.(dog(x) & disappear(x))

In addition to the existential quantifier, first-order logic offers us the **universal quantifier**  $\forall x$  (“for all *x*”), illustrated in **Example 10-20**.

*Example 10-20.*

1.  $\forall x.(dog(x) \rightarrow disappear(x))$
2. Everything has the property that if it is a dog, it disappears.
3. Every dog disappeared.

Here is the NLTK counterpart of **a** :

*Example 10-21.*

all x.(dog(x) -> disappear(x))

Although **a** is the standard first-order logic translation of **c** , the truth conditions aren’t necessarily what you expect. The formula says that *if* some *x* is a dog, then *x* disappears—but it doesn’t say that there are any dogs. So in a situation where there are no dogs, **a** will still come out true. (Remember that  $(P \rightarrow Q)$  is true when *P* is false.) Now you might argue that *every dog disappeared* does presuppose the existence of dogs, and that the logic formalization is simply wrong. But it is possible to find other examples that lack



every occurrence of 'ate' in *astring* by '8', even though there may in fact be no such occurrences (Table 3-2).

We have seen a number of examples where variables are bound by quantifiers. What happens in formulas such as the following?

```
((exists x. dog(x)) -> bark(x))
```

The scope of the `exists x` quantifier is `dog(x)`, so the occurrence of `x` in `bark(x)` is unbound. Consequently it can become bound by some other quantifier, for example, `all x` in the next formula:

```
all x. ((exists x. dog(x)) -> bark(x))
```

In general, an occurrence of a variable `x` in a formula  $\varphi$  is **free** in  $\varphi$  if that occurrence doesn't fall within the scope of `all x` or `some x` in  $\varphi$ . Conversely, if `x` is free in formula  $\varphi$ , then it is **bound** in `all x.  $\varphi$`  and `exists x.  $\varphi$` . If all variable occurrences in a formula are bound, the formula is said to be **closed**.

We mentioned before that the `parse()` method of NLTK's `LogicParser` returns objects of class `Expression`. Each instance `expr` of this class comes with a method `free()`, which returns the set of variables that are free in `expr`.

```
>>> lp = nltk.LogicParser()
>>> lp.parse('dog(cyril)').free()
set([])
>>> lp.parse('dog(x)').free()
set([Variable('x')])
>>> lp.parse('own(angus, cyril)').free()
set([])
>>> lp.parse('exists x.dog(x)').free()
set([])
>>> lp.parse('((some x. walk(x)) -> sing(x))').free()
set([Variable('x')])
>>> lp.parse('exists x.own(y, x)').free()
set([Variable('y')])
```

FIRST-ORDER THEOREM PROVING

Recall the constraint on *to the north of*, which we proposed earlier as

Example 10-10:

Example 10-22.

if *x* is to the north of *y* then *y* is not to the north of *x*.

We observed that propositional logic is not expressive enough to represent generalizations about binary predicates, and as a result we did not properly capture the argument *Sylvania is to the north of Freedonia. Therefore, Freedonia is not to the north of Sylvania*.

You have no doubt realized that first-order logic, by contrast, is ideal for formalizing such rules:

```
all x. all y. (north_of(x, y) -> -north_of(y, x))
```

Even better, we can perform automated inference to show the validity of the argument.

The general case in theorem proving is to determine whether a formula that we want to prove (a **proof goal**) can be derived by a finite sequence of inference steps from a list of assumed formulas. We write this as  $A \vdash g$ , where  $A$  is a (possibly empty) list of assumptions, and  $g$  is a proof goal. We will illustrate this with NLTK's interface to the theorem prover Prover9. First, we parse the



```
>>> NotFnS = lp.parse('-north_of(f, s)') ❶
>>> SnF = lp.parse('north_of(s, f)') ❷
>>> R = lp.parse('all x. all y. (north_of(x, y) -> -north_of(y, x)')
>>> prover = nltk.Prover9() ❸
>>> prover.prove(NotFnS, [SnF, R]) ❹
True ❺
```

Happily, the theorem prover agrees with us that the argument is valid. By contrast, it concludes that it is not possible to infer `north_of(f, s)` from our assumptions:

```
>>> FnS = lp.parse('north_of(f, s)')
>>> prover.prove(FnS, [SnF, R])
False
```

### SUMMARIZING THE LANGUAGE OF FIRST-ORDER LOGIC

We'll take this opportunity to restate our earlier syntactic rules for propositional logic and add the formation rules for quantifiers; together, these give us the syntax of first-order logic. In addition, we make explicit the types of the expressions involved. We'll adopt the convention that  $\langle e, t \rangle$  is the type of a predicate that combines with  $n$  arguments of type  $e$  to yield an expression of type  $t$ . In this case, we say that  $n$  is the **arity** of the predicate.

1. If  $P$  is a predicate of type  $\langle e, t \rangle$ , and  $\alpha_1, \dots, \alpha_n$  are terms of type  $e$ , then  $P(\alpha_1, \dots, \alpha_n)$  is of type  $t$ .
2. If  $\alpha$  and  $\beta$  are both of type  $e$ , then  $(\alpha = \beta)$  and  $(\alpha \neq \beta)$  are of type  $t$ .
3. If  $\varphi$  is of type  $t$ , then so is  $\neg \varphi$ .
4. If  $\varphi$  and  $\psi$  are of type  $t$ , then so are  $(\varphi \ \& \ \psi)$ ,  $(\varphi \ \mid \ \psi)$ ,  $(\varphi \ \rightarrow \ \psi)$ , and  $(\varphi \ \leftrightarrow \ \psi)$ .
5. If  $\varphi$  is of type  $t$ , and  $x$  is a variable of type  $e$ , then  $\text{exists } x. \varphi$  and  $\text{all } x. \varphi$  are of type  $t$ .

**Table 10-3** summarizes the new logical constants of the `logic` module, and two of the methods of `Expressions`.

Table 10-3. Summary of new logical relations and operators required for first-order logic

Example	Description
=	Equality
!=	Inequality
exists	Existential quantifier
all	Universal quantifier



## TRUTH IN MODEL

We have looked at the syntax of first-order logic, and in [The Semantics of English Sentences](#) we will examine the task of translating English into first-order logic. Yet as we argued in [Natural Language Understanding](#), this gets us further forward only if we can give a meaning to sentences of first-order logic. In other words, we need to give a *truth-conditional semantics* to first-order logic. From the point of view of computational semantics, there are obvious limits to how far one can push this approach. Although we want to talk about sentences being true or false in situations, we only have the means of representing situations in the computer in a symbolic manner. Despite this limitation, it is still possible to gain a clearer picture of truth-conditional semantics by encoding models in NLTK.

Given a first-order logic language  $L$ , a model  $M$  for  $L$  is a pair  $\langle D, Val \rangle$ , where  $D$  is a non-empty set called the **domain** of the model, and  $Val$  is a function called the **valuation function**, which assigns values from  $D$  to expressions of  $L$  as follows:

1. For every individual constant  $c$  in  $L$ ,  $Val(c)$  is an element of  $D$ .
2. For every predicate symbol  $P$  of arity  $n \geq 0$ ,  $Val(P)$  is a function from  $D^n$  to  $\{True, False\}$ . (If the arity of  $P$  is 0, then  $Val(P)$  is simply a truth value, and  $P$  is regarded as a propositional symbol.)

According to 2, if  $P$  is of arity 2, then  $Val(P)$  will be a function  $f$  from pairs of elements of  $D$  to  $\{True, False\}$ . In the models we shall build in NLTK, we'll adopt a more convenient alternative, in which  $Val(P)$  is a set  $S$  of pairs, defined as follows:

*Example 10-23.*

$S = \{s \mid f(s) = True\}$

Such an  $f$  is called the **characteristic function** of  $S$  (as discussed in the further readings).

Relations are represented semantically in NLTK in the standard set-theoretic way: as sets of tuples. For example, let's suppose we have a domain of discourse consisting of the individuals Bertie, Olive, and Cyril, where Bertie is a boy, Olive is a girl, and Cyril is a dog. For mnemonic reasons, we use `b`, `o`, and `c` as the corresponding labels in the model. We can declare the domain as follows:

```
>>> dom = set(['b', 'o', 'c'])
```

We will use the utility function `parse_valuation()` to convert a sequence of strings of the form *symbol => value* into a `Valuation` object.

```
>>> v = """
... bertie => b
... olive => o
... cyril => c
... boy => (b)
... girl => (o)
... dog => (c)
... walk => (o, c)
... see => ((b, o), (c, b), (o, c))
... """
>>> val = nltk.parse_valuation(v)
>>> print val
{'bertie': 'b',
 'boy': set([('b',)]),
 'cyril': 'c',
 'dog': set([('c',)]),
 'girl': set([('o',)]),
 'olive': 'o',
 'see': set([('o', 'c'), ('c', 'b'), ('b', 'o')]),
 'walk': set([('c',), ('o',)])}
```

So according to this valuation, the value of `see` is a set of tuples such that Bertie





## NOTE

**Your Turn:** Draw a picture of the domain *dom* and the sets corresponding to each of the unary predicates, by analogy with the diagram shown in [Figure 10-2](#).

You may have noticed that our unary predicates (i.e. `boy`, `girl`, `dog`) also come out as sets of singleton tuples, rather than just sets of individuals. This is a convenience which allows us to have a uniform treatment of relations of any arity. A predication of the form  $P(\tau_1, \dots, \tau_n)$ , where  $P$  is of arity  $n$ , comes out true just in case the tuple of values corresponding to  $(\tau_1, \dots, \tau_n)$  belongs to the set of tuples in the value of  $P$ .

```
>>> ('o', 'c') in val['see']
True
>>> ('b',) in val['boy']
True
```

## INDIVIDUAL VARIABLES AND ASSIGNMENTS

In our models, the counterpart of a context of use is a variable **assignment**. This is a mapping from individual variables to entities in the domain. Assignments are created using the `Assignment` constructor, which also takes the model's domain of discourse as a parameter. We are not required to actually enter any bindings, but if we do, they are in a *(variable, value)* format similar to what we saw earlier for valuations.

```
>>> g = nltk.Assignment(dom, [('x', 'o'), ('y', 'c')])
>>> g
{'y': 'c', 'x': 'o'}
```

In addition, there is a `print()` format for assignments which uses a notation closer to that often found in logic textbooks:

```
>>> print g
g[c/y][o/x]
```

Let's now look at how we can evaluate an atomic formula of first-order logic. First, we create a model, and then we call the `evaluate()` method to compute the truth value:

```
>>> m = nltk.Model(dom, val)
>>> m.evaluate('see(olive, y)', g)
True
```

What's happening here? We are evaluating a formula which is similar to our earlier example, `see(olive, cyril)`. However, when the interpretation function encounters the variable `y`, rather than checking for a value in `val`, it asks the variable assignment `g` to come up with a value:

```
>>> g['y']
'c'
```

Since we already know that individuals `o` and `c` stand in the *see* relation, the value `True` is what we expected. In this case, we can say that assignment `g` **satisfies** the formula `see(olive, y)`. By contrast, the following formula evaluates to `False` relative to `g` (check that you see why this is).

```
>>> m.evaluate('see(y, x)', g)
```



In our approach (though not in standard first-order logic), variable assignments are *partial*. For example, `g` says nothing about any variables apart from `x` and `y`. The method `purge()` clears all bindings from an assignment.

```
>>> g.purge()
>>> g
{}

```

If we now try to evaluate a formula such as `see(olive, y)` relative to `g`, it is like trying to interpret a sentence containing a *him* when we don't know what *him* refers to. In this case, the evaluation function fails to deliver a truth value.

```
>>> m.evaluate('see(olive, y)', g)
'Undefined'

```

Since our models already contain rules for interpreting Boolean operators, arbitrarily complex formulas can be composed and evaluated.

```
>>> m.evaluate('see(bertie, olive) & boy(bertie) & ~walk(bertie)'
True

```

The general process of determining truth or falsity of a formula in a model is called **model checking**.

## QUANTIFICATION

One of the crucial insights of modern logic is that the notion of variable satisfaction can be used to provide an interpretation for quantified formulas. Let's use [Example 10-24](#) as an example.

*Example 10-24.*

```
exists x.(girl(x) & walk(x))

```

When is it true? Let's think about all the individuals in our domain, i.e., in `dom`. We want to check whether any of these individuals has the property of being a girl and walking. In other words, we want to know if there is some *u* in `dom` such that `g[u/x]` satisfies the open formula [Example 10-25](#).

*Example 10-25.*

```
girl(x) & walk(x)

```

Consider the following:

```
>>> m.evaluate('exists x.(girl(x) & walk(x))', g)
True

```

`evaluate()` returns `True` here because there is some *u* in `dom` such that [Example 10-25](#) is satisfied by an assignment which binds `x` to *u*. In fact, `o` is such a *u*:

```
>>> m.evaluate('girl(x) & walk(x)', g.add('x', 'o'))
True

```

One useful tool offered by NLTK is the `satisfiers()` method. This returns a set of all the individuals that satisfy an open formula. The method parameters are a parsed formula, a variable, and an assignment. Here are a few examples:

```
>>> fmlal = lp.parse('girl(x) | boy(x)')
>>> m.satisfiers(fmlal, 'x', g)
set(['b', 'o'])

```



```
>>> m.satisfiers(fmla3, 'x', g)
set(['b', 'o'])
```

It's useful to think about why `fmla2` and `fmla3` receive the values they do. The truth conditions for  $\rightarrow$  mean that `fmla2` is equivalent to  $\neg \text{girl}(x) \mid \text{walk}(x)$ , which is satisfied by something that either isn't a girl or walks. Since neither `b` (Bertie) nor `c` (Cyril) are girls, according to model `m`, they both satisfy the whole formula. And of course `o` satisfies the formula because `o` satisfies both disjuncts. Now, since every member of the domain of discourse satisfies `fmla2`, the corresponding universally quantified formula is also true.

```
>>> m.evaluate('all x.(girl(x) -> walk(x))', g)
True
```

In other words, a universally quantified formula  $\forall x.\phi$  is true with respect to `g` just in case for every `u`,  $\phi$  is true with respect to `g[u/x]`.

#### NOTE

**Your Turn:** Try to figure out, first with pencil and paper, and then using `m.evaluate()`, what the truth values are for `all x.(girl(x) & walk(x))` and `exists x.(boy(x) -> walk(x))`. Make sure you understand why they receive these values.

### QUANTIFIER SCOPE AMBIGUITY

What happens when we want to give a formal representation of a sentence with *two* quantifiers, such as the following?

*Example 10-26.*

Everybody admires someone.

There are (at least) two ways of expressing [Example 10-26](#) in first-order logic:

*Example 10-27.*

1. `all x.(person(x) -> exists y.(person(y) & admire(x,y)))`
2. `exists y.(person(y) & all x.(person(x) -> admire(x,y)))`

Can we use both of these? The answer is yes, but they have different meanings.

`b` is logically stronger than `a` : it claims that there is a unique person, say,

Bruce, who is admired by everyone. `a` , on the other hand, just requires that for every person `u`, we can find some person `u'` whom `u` admires; but this could be a different person `u'` in each case. We distinguish between `a` and `b` in terms of the **scope** of the quantifiers. In the first,  $\forall$  has wider scope than  $\exists$ , whereas in

`b` , the scope ordering is reversed. So now we have two ways of representing


the meaning of [Example 10-26](#), and they are both quite legitimate. In other words, we are claiming that [Example 10-26](#) is *ambiguous* with respect to quantifier scope, and the formulas in [Example 10-27](#) give us a way to make the two readings explicit. However, we are not just interested in associating two distinct representations with [Example 10-26](#); we also want to show in detail how the two representations lead to different conditions for truth in a model.



```
>>> v2 = """
... bruce => b
... cyril => c
... elspeth => e
... julia => j
... matthew => m
... person => {b, e, j, m}
... admire => {(j, b), (b, b), (m, e), (e, m), (c, a)}
... """
>>> val2 = nltk.parse_valuation(v2)
```

The *admire* relation can be visualized using the mapping diagram shown in [Example 10-28](#).

*Example 10-28.*



In [Example 10-28](#), an arrow between two individuals  $x$  and  $y$  indicates that  $x$  admires  $y$ . So  $j$  and  $b$  both admire  $b$  (Bruce is very vain), while  $e$  admires  $m$  and  $m$  admires  $e$ . In this model, formula **a** is true but **b** is false. One way of exploring these results is by using the `satisfiers()` method of `Model` objects.

```
>>> dom2 = val2.domain
>>> m2 = nltk.Model(dom2, val2)
>>> g2 = nltk.Assignment(dom2)
>>> fmla4 = lp.parse('(person(x) -> exists y.(person(y) & admire(x, y)))')
>>> m2.satisfiers(fmla4, 'x', g2)
set(['a', 'c', 'b', 'e', 'j', 'm'])
```

This shows that `fmla4` holds of every individual in the domain. By contrast, consider the formula `fmla5`; this has no satisfiers for the variable  $y$ .

```
>>> fmla5 = lp.parse('(person(y) & all x.(person(x) -> admire(x, y)))')
>>> m2.satisfiers(fmla5, 'y', g2)
set([])
```

That is, there is no person that is admired by everybody. Taking a different open formula, `fmla6`, we can verify that there is a person, namely Bruce, who is admired by both Julia and Bruce.

```
>>> fmla6 = lp.parse('(person(y) & all x.((x = bruce | x = julia) -> admire(x, y)))')
>>> m2.satisfiers(fmla6, 'y', g2)
set(['b', 'j'])
```



## NOTE

**Your Turn:** Devise a new model based on  $\mathfrak{M}_2$  such that **a** comes out false in your model; similarly, devise a new model such that **b** comes out true.

## MODEL BUILDING

We have been assuming that we already had a model, and wanted to check the truth of a sentence in the model. By contrast, model building tries to create a new model, given some set of sentences. If it succeeds, then we know that the set is consistent, since we have an existence proof of the model.

We invoke the Mace4 model builder by creating an instance of `Mace()` and calling its `build_model()` method, in an analogous way to calling the Prover9 theorem prover. One option is to treat our candidate set of sentences as assumptions, while leaving the goal unspecified. The following interaction shows how both  $[a, c1]$  and  $[a, c2]$  are consistent lists, since Mace succeeds in building a model for each of them, whereas  $[c1, c2]$  is inconsistent.

```
>>> a3 = lp.parse('exists x. (man(x) & walks(x))')
>>> c1 = lp.parse('mortal(socrates)')
>>> c2 = lp.parse('!mortal(socrates)')
>>> mb = nltk.Mace(5)
>>> print mb.build_model(None, [a3, c1])
True
>>> print mb.build_model(None, [a3, c2])
True
>>> print mb.build_model(None, [c1, c2])
False
```

We can also use the model builder as an adjunct to the theorem prover. Let's suppose we are trying to prove  $A \vdash g$ , i.e., that  $g$  is logically derivable from assumptions  $A = [a1, a2, \dots, an]$ . We can feed this same input to Mace4, and the model builder will try to find a counterexample, that is, to show that  $g$  does *not* follow from  $A$ . So, given this input, Mace4 will try to find a model for the assumptions  $A$  together with the negation of  $g$ , namely the list  $A' = [a1, a2, \dots, an, \neg g]$ . If  $g$  fails to follow from  $S$ , then Mace4 may well return with a counterexample faster than Prover9 concludes that it cannot find the required proof. Conversely, if  $g$  is provable from  $S$ , Mace4 may take a long time unsuccessfully trying to find a countermodel, and will eventually give up.

Let's consider a concrete scenario. Our assumptions are the list *[There is a woman that every man loves, Adam is a man, Eve is a woman]*. Our conclusion is *Adam loves Eve*. Can Mace4 find a model in which the premises are true but the conclusion is false? In the following code, we use `MaceCommand()`, which will let us inspect the model that has been built.

```
>>> a4 = lp.parse('exists y. (woman(y) & all x. (man(x) -> love(x, y)))')
>>> a5 = lp.parse('man(adam)')
>>> a6 = lp.parse('woman(eve)')
>>> g = lp.parse('!love(adam, eve)')
>>> mc = nltk.MaceCommand(g, assumptions=[a4, a5, a6])
>>> mc.build_model()
True
```

So the answer is yes: Mace4 found a countermodel in which there is some woman other than Eve that Adam loves. But let's have a closer look at Mace4's model, converted to the format we use for valuations:



```
'love': set([('a', 'b')]),  
'man': set([('a',)]),  
'woman': set([('a',), ('b',)])}
```

The general form of this valuation should be familiar to you: it contains some individual constants and predicates, each with an appropriate kind of value. What might be puzzling is the **C1**. This is a “Skolem constant” that the model builder introduces as a representative of the existential quantifier. That is, when the model builder encountered the **exists y** part of **a4**, it knew that there is some individual **b** in the domain which satisfies the open formula in the body of **a4**. However, it doesn’t know whether **b** is also the denotation of an individual constant anywhere else in its input, so it makes up a new name for **b** on the fly, namely **C1**. Now, since our premises said nothing about the individual constants **adam** and **eve**, the model builder has decided there is no reason to treat them as denoting different entities, and they both get mapped to **a**. Moreover, we didn’t specify that **man** and **woman** denote disjoint sets, so the model builder lets their denotations overlap. This illustrates quite dramatically the implicit knowledge that we bring to bear in interpreting our scenario, but which the model builder knows nothing about. So let’s add a new assumption which makes the sets of men and women disjoint. The model builder still produces a countermodel, but this time it is more in accord with our intuitions about the situation:

```
>>> a7 = lp.parse('all x. (man(x) -> ~woman(x))')  
>>> g = lp.parse('love(adam,eve)')  
>>> mc = nltk.MakeCommand(g, assumptions=[a4, a5, a6, a7])  
>>> mc.build_model()  
True  
>>> print mc.valuation  
{ 'C1': 'c',  
  'adam': 'a',  
  'eve': 'b',  
  'love': set([('a', 'c')]),  
  'man': set([('a',)]),  
  'woman': set([('b',), ('c',)])}
```

On reflection, we can see that there is nothing in our premises which says that Eve is the only woman in the domain of discourse, so the countermodel in fact is acceptable. If we wanted to rule it out, we would have to add a further assumption such as **exists y. all x. (woman(x) -> (x = y))** to ensure that there is only one woman in the model.

## The Semantics of English Sentences

### COMPOSITIONAL SEMANTICS IN FEATURE-BASED GRAMMAR

At the beginning of the chapter we briefly illustrated a method of building semantic representations on the basis of a syntactic parse, using the grammar framework developed in [Chapter 9](#). This time, rather than constructing an SQL query, we will build a logical form. One of our guiding ideas for designing such grammars is the **Principle of Compositionality**. (Also known as Frege’s Principle; see [Partee, 1995] for the formulation given.)


**Principle of Compositionality:** the meaning of a whole is a function of the meanings of the parts and of the way they are syntactically combined.

We will assume that the semantically relevant parts of a complex expression are given by a theory of syntactic analysis. Within this chapter, we will take it for granted that expressions are parsed against a context-free grammar. However, this is not entailed by the Principle of Compositionality.

Our goal now is to integrate the construction of a semantic representation in a manner that can be smoothly with the process of parsing. [Example 10-29](#) illustrates a first approximation to the kind of analyses we would like to build.

*Example 10-29*





In [Example 10-29](#), the SEM value at the root node shows a semantic representation for the whole sentence, while the SEM values at lower nodes show semantic representations for constituents of the sentence. Since the values of SEM have to be treated in a special manner, they are distinguished from other feature values by being enclosed in angle brackets.

So far, so good, but how do we write grammar rules that will give us this kind of result? Our approach will be similar to that adopted for the grammar *sql0.fcfg* at the start of this chapter, in that we will assign semantic representations to lexical nodes, and then compose the semantic representations for each phrase from those of its child nodes. However, in the present case we will use function application rather than string concatenation as the mode of composition. To be more specific, suppose we have NP and VP constituents with appropriate values for their SEM nodes. Then the SEM value of an S is handled by a rule like [Example 10-30](#). (Observe that in the case where the value of SEM is a variable, we omit the angle brackets.)

*Example 10-30.*

---

S[SEM=<?vp(?np)>] -> NP[SEM=?np] VP[SEM=?vp]

[Example 10-30](#) tells us that given some SEM value ?np for the subject NP and some SEM value ?vp for the VP, the SEM value of the S parent is constructed by applying ?vp as a function expression to ?np. From this, we can conclude that ?vp has to denote a function which has the denotation of ?np in its domain. [Example 10-30](#) is a nice example of building semantics using the principle of compositionality.

To complete the grammar is very straightforward; all we require are the rules shown here:

---

```

VP[SEM=?v] -> IV[SEM=?v]
NP[SEM=<cyrl>] -> 'Cyril'
IV[SEM=<\x.bark(x)>] -> 'barks'
  
```

---

The VP rule says that the parent's semantics is the same as the head child's semantics. The two lexical rules provide non-logical constants to serve as the semantic values of *Cyril* and *barks* respectively. There is an additional piece of notation in the entry for *barks* which we will explain shortly.

Before launching into compositional semantic rules in more detail, we need to add a new tool to our kit, namely the  $\lambda$ -calculus. This provides us with an invaluable tool for combining expressions of first-order logic as we assemble a meaning representation for an English sentence.

## THE $\lambda$ -CALCULUS

In [Computing with Language: Simple Statistics](#), we pointed out that mathematical set notation was a helpful method of specifying properties *P* of words that we wanted to select from a document. We illustrated this with [Example 10-31](#), which we glossed as “the set of all *w* such that *w* is an element of



$\{w \mid w \in V \ \& \ P(w)\}$

It turns out to be extremely useful to add something to first-order logic that will achieve the same effect. We do this with the  **$\lambda$ -operator** (pronounced “lambda”). The  $\lambda$  counterpart to [Example 10-31](#) is [Example 10-32](#). (Since we are not trying to do set theory here, we just treat  $V$  as a unary predicate.)

*Example 10-32.*

$\lambda w. (V(w) \ \& \ P(w))$

#### NOTE

$\lambda$  expressions were originally designed by Alonzo Church to represent computable functions and to provide a foundation for mathematics and logic. The theory in which  $\lambda$  expressions are studied is known as the  $\lambda$ -calculus. Note that the  $\lambda$ -calculus is not part of first-order logic—both can be used independently of the other.

$\lambda$  is a binding operator, just as the first-order logic quantifiers are. If we have an open formula, such as [a](#) , then we can bind the variable  $x$  with the  $\lambda$  operator, as shown in [b](#) . The corresponding NLTK representation is given in [c](#) .

*Example 10-33.*

1.  $(walk(x) \ \& \ chew\_gum(x))$
2.  $\lambda x.(walk(x) \ \& \ chew\_gum(x))$
3. `\x.(walk(x) \ \& \ chew_gum(x))`

Remember that `\` is a special character in Python strings. We must either escape it (with another `\`), or else use “raw strings” ([Regular Expressions for Detecting Word Patterns](#)) as shown here:

```
>>> lp = nltk.LogicParser()
>>> e = lp.parse(r'\x.(walk(x) \& chew_gum(x))')
>>> e
<LambdaExpression \x.(walk(x) \& chew_gum(x))>
>>> e.free()
set([])
>>> print lp.parse(r'\x.(walk(x) \& chew_gum(y))')
\x.(walk(x) \& chew_gum(y))
```

We have a special name for the result of binding the variables in an expression:  **$\lambda$ -abstraction**. When you first encounter  $\lambda$ -abstracts, it can be hard to get an

intuitive sense of their meaning. A couple of English glosses for [b](#) are: “be an  $x$  such that  $x$  walks and  $x$  chews gum” or “have the property of walking and chewing gum.” It has often been suggested that  $\lambda$ -abstracts are good representations for verb phrases (or subjectless clauses), particularly when these occur as arguments in their own right. This is illustrated in [a](#) and its

translation, [b](#) .

*Example 10-34.*

1. To walk and chew gum is hard





So the general picture is this: given an open formula  $\phi$  with free variable  $x$ , abstracting over  $x$  yields a property expression  $\lambda x.\phi$ —the property of being an  $x$  such that  $\phi$ . Here’s a more official version of how abstracts are built:

*Example 10-35.*

If  $\alpha$  is of type  $\tau$ , and  $x$  is a variable of type  $e$ , then  $\lambda x.\alpha$  is of type  $\langle e, \tau \rangle$ .

**b** illustrated a case where we say something about a property, namely that it is hard. But what we usually do with properties is attribute them to individuals. And in fact, if  $\phi$  is an open formula, then the abstract  $\lambda x.\phi$  can be used as a unary predicate. In [Example 10-36](#), **b** is predicated of the term *gerald*.

*Example 10-36.*

$\lambda x.(\text{walk}(x) \ \& \ \text{chew\_gum}(x)) \ (\text{gerald})$

Now [Example 10-36](#) says that Gerald has the property of walking and chewing gum, which has the same meaning as [Example 10-37](#).

*Example 10-37.*

$(\text{walk}(\text{gerald}) \ \& \ \text{chew\_gum}(\text{gerald}))$

What we have done here is remove the  $\lambda x$  from the beginning of  $\lambda x.(\text{walk}(x) \ \& \ \text{chew\_gum}(x))$  and replaced all occurrences of  $x$  in  $(\text{walk}(x) \ \& \ \text{chew\_gum}(x))$  by *gerald*. We’ll use  $\alpha[\beta/x]$  as notation for the operation of replacing all free occurrences of  $x$  in  $\alpha$  by the expression  $\beta$ . So

$(\text{walk}(x) \ \& \ \text{chew\_gum}(x)) [\text{gerald}/x]$

represents the same expression as [Example 10-37](#). The “reduction” of [Example 10-36](#) to [Example 10-37](#) is an extremely useful operation in simplifying semantic representations, and we shall use it a lot in the rest of this chapter. The operation is often called  **$\beta$ -reduction**. In order for it to be semantically justified, we want it to hold that  $\lambda x. \alpha(\beta)$  has the same semantic value as  $\alpha[\beta/x]$ . This is indeed true, subject to a slight complication that we will come to shortly. In order to carry out  $\beta$ -reduction of expressions in NLTK, we can call the `simplify()` method [1](#).

```
>>> e = lp.parse(r'\x.(\text{walk}(x) \ \& \ \text{chew\_gum}(x)) (\text{gerald})')
>>> print e
\x.(\text{walk}(x) \ \& \ \text{chew\_gum}(x)) (\text{gerald})
>>> print e.simplify() 1
(\text{walk}(\text{gerald}) \ \& \ \text{chew\_gum}(\text{gerald}))
```

Although we have so far only considered cases where the body of the  $\lambda$ -abstract is an open formula, i.e., of type  $t$ , this is not a necessary restriction; the body can be any well-formed expression. Here’s an example with two  $\lambda$ s:

*Example 10-38.*

$\lambda x.\lambda y.(\text{dog}(x) \ \& \ \text{own}(y, \ x))$

Just as **b** plays the role of a unary predicate, [Example 10-38](#) works like a binary predicate: it can be applied directly to two arguments [1](#). The `LogicParser` allows nested  $\lambda$ s such as  $\lambda x.\lambda y.$  to be written in the abbreviated form  $\lambda x \ y.$  [1](#).

```
>>> print lp.parse(r'\x.\y.(\text{dog}(x) \ \& \ \text{own}(y, \ x)) (\text{cyril})').simplify
\y.(\text{dog}(\text{cyril}) \ \& \ \text{own}(y, \ \text{cyril}))
```



All our  $\lambda$ -abstracts so far have involved the familiar first-order variables:  $x$ ,  $y$ , and so on—variables of type  $e$ . But suppose we want to treat one abstract, say,  $\backslash x.\text{walk}(x)$ , as the *argument* of another  $\lambda$ -abstract? We might try this:

```
\y.y(angus) (\x.walk(x))
```

But since the variable  $y$  is stipulated to be of type  $e$ ,  $\backslash y.y(\text{angus})$  only applies to arguments of type  $e$  while  $\backslash x.\text{walk}(x)$  is of type  $(e, t)$ ! Instead, we need to allow abstraction over variables of higher type. Let's use  $P$  and  $Q$  as variables of type  $(e, t)$ , and then we can have an abstract such as  $\backslash P.P(\text{angus})$ . Since  $P$  is of type  $(e, t)$ , the whole abstract is of type  $((e, t), t)$ . Then  $\backslash P.P(\text{angus}) (\backslash x.\text{walk}(x))$  is legal, and can be simplified via  $\beta$ -reduction to  $\backslash x.\text{walk}(x)(\text{angus})$  and then again to  $\text{walk}(\text{angus})$ .

When carrying out  $\beta$ -reduction, some care has to be taken with variables.

Consider, for example, the  $\lambda$ -terms  $a$  and  $b$ , which differ only in the identity of a free variable.

*Example 10-39.*

- 1.  $\backslash y.\text{see}(y, x)$
- 2.  $\backslash y.\text{see}(y, z)$

Suppose now that we apply the  $\lambda$ -term  $\backslash P.\text{exists } x.P(x)$  to each of these terms:

*Example 10-40.*

- 1.  $\backslash P.\text{exists } x.P(x)(\backslash y.\text{see}(y, x))$
- 2.  $\backslash P.\text{exists } x.P(x)(\backslash y.\text{see}(y, z))$

We pointed out earlier that the results of the application should be semantically equivalent. But if we let the free variable  $x$  in  $a$  fall inside the scope of the existential quantifier in  $a$ , then after reduction, the results will be different:

*Example 10-41.*

- 1.  $\text{exists } x.\text{see}(x, x)$
- 2.  $\text{exists } x.\text{see}(x, z)$

$a$  means there is some  $x$  that sees him/herself, whereas  $b$  means that there is some  $x$  that sees an unspecified individual  $z$ . What has gone wrong here? Clearly, we want to forbid the kind of variable “capture” shown in  $a$ .

In order to deal with this problem, let's step back a moment. Does it matter what particular name we use for the variable bound by the existential quantifier in the function expression of  $a$ ? The answer is no. In fact, given any variable-binding expression (involving  $\forall, \exists$ , or  $\lambda$ ), the name chosen for the bound variable is completely arbitrary. For example,  $\text{exists } x.P(x)$  and  $\text{exists } y.P(y)$  are equivalent; they are called  $\alpha$ -equivalents, or **alphabetic variants**. The process of relabeling bound variables is known as  **$\alpha$ -conversion**. When we test for equality of `VariableBinderExpressions` in the `logic` module (i.e., using `==`), we are in fact testing for  $\alpha$ -equivalence:

```
>>> e1 = lp.parse('exists x.P(x)')
```



```
exists z.P(z)
>>> e1 == e2
True
```

When  $\beta$ -reduction is carried out on an application  $f(a)$ , we check whether there are free variables in  $a$  that also occur as bound variables in any subterms of  $f$ . Suppose, as in the example just discussed, that  $x$  is free in  $a$ , and that  $f$  contains the subterm  $\text{exists } x.P(x)$ . In this case, we produce an alphabetic variant of  $\text{exists } x.P(x)$ , say,  $\text{exists } z1.P(z1)$ , and then carry on with the reduction. This relabeling is carried out automatically by the  $\beta$ -reduction code in `logic`, and the results can be seen in the following example:

```
>>> e3 = lp.parse('\P.exists x.P(x) (\y.see (y, x))')
>>> print e3
(\P.exists x.P(x)) (\y.see (y,x))
>>> print e3.simplify()
exists z1.see(z1,x)
```

#### NOTE

As you work through examples like these in the following sections, you may find that the logical expressions which are returned have different variable names; for example, you might see `z14` in place of `z1` in the preceding formula. This change in labeling is innocuous—in fact, it is just an illustration of alphabetic variants.

After this excursus, let's return to the task of building logical forms for English sentences.

### QUANTIFIED NPS

At the start of this section, we briefly described how to build a semantic representation for *Cyril barks*. You would be forgiven for thinking this was all too easy—surely there is a bit more to building compositional semantics. What about quantifiers, for instance? Right, this is a crucial issue. For example, we want `a` to be given the logical form in `b`. How can this be accomplished?

*Example 10-42.*

1. A dog barks.
2. `exists x.(dog(x) & bark(x))`

Let's make the assumption that our *only* operation for building complex semantic representations is function application. Then our problem is this: how do we give a semantic representation to the quantified NPs *a dog* so that it can be combined with *bark* to give the result in `b`? As a first step, let's make the subject's SEM value act as the function expression rather than the argument. (This is sometimes called **type-raising**.) Now we are looking for a way of instantiating `?np` so that `[SEM=<?np(\x.bark(x))>]` is equivalent to `[SEM=<exists x.(dog(x) & bark(x))>]`. Doesn't this look a bit reminiscent of carrying out  $\beta$ -reduction in the  $\lambda$ -calculus? In other words, we want a  $\lambda$ -term  $M$  to replace `?np` so that applying  $M$  to `\x.bark(x)` yields `b`. To do this, we replace the occurrence of

`\x.bark(x)` in `b` by a predicate variable  $P$ , and bind the variable with  $\lambda$ , as shown in **Example 10-43**.



$\lambda P. \text{exists } x. (\text{dog}(x) \ \& \ P(x))$

We have used a different style of variable in [Example 10-43](#)—that is, 'P' rather than 'x' or 'y'—to signal that we are abstracting over a different kind of object—not an individual, but a function expression of type  $\langle e, t \rangle$ . So the type of [Example 10-43](#) as a whole is  $\langle \langle e, t \rangle, t \rangle$ . We will take this to be the type of NPs in general. To illustrate further, a universally quantified NP will look like [Example 10-44](#).

*Example 10-44.*

$\lambda P. \text{all } x. (\text{dog}(x) \rightarrow P(x))$

We are pretty much done now, except that we also want to carry out a further abstraction plus application for the process of combining the semantics of the determiner *a*, namely [Example 10-45](#), with the semantics of *dog*.

*Example 10-45.*

$\lambda Q \lambda P. \text{exists } x. (Q(x) \ \& \ P(x))$

Applying [Example 10-45](#) as a function expression to  $\lambda x. \text{dog}(x)$  yields [Example 10-43](#), and applying that to  $\lambda x. \text{bark}(x)$  gives us  $\lambda P. \text{exists } x. (\text{dog}(x) \ \& \ P(x)) (\lambda x. \text{bark}(x))$ . Finally, carrying out  $\beta$ -reduction yields just what we wanted, namely **b**.

## TRANSITIVE VERBS

Our next challenge is to deal with sentences containing transitive verbs, such as [Example 10-46](#).

*Example 10-46.*

Angus chases a dog.

The output semantics that we want to build is  $\text{exists } x. (\text{dog}(x) \ \& \ \text{chase}(\text{angus}, x))$ . Let's look at how we can use  $\lambda$ -abstraction to get this result. A significant constraint on possible solutions is to require that the semantic representation of *a dog* be independent of whether the NP acts as subject or object of the sentence. In other words, we want to get the formula just shown as our output while sticking to [Example 10-43](#) as the NP semantics. A second constraint is that VPs should have a uniform type of interpretation, regardless of whether they consist of just an intransitive verb or a transitive verb plus object. More specifically, we stipulate that VPs are always of type  $\langle e, t \rangle$ . Given these constraints, here's a semantic representation for *chases a dog* that does the trick.

*Example 10-47.*

$\lambda y. \text{exists } x. (\text{dog}(x) \ \& \ \text{chase}(y, x))$

Think of [Example 10-47](#) as the property of being a *y* such that for some dog *x*, *y* chases *x*; or more colloquially, being a *y* who chases a dog. Our task now resolves to designing a semantic representation for *chases* which can combine with [Example 10-43](#) so as to allow [Example 10-47](#) to be derived.

Let's carry out the inverse of  $\beta$ -reduction on [Example 10-47](#), giving rise to [Example 10-48](#).

*Example 10-48.*

$\lambda P. \text{exists } x. (\text{dog}(x) \ \& \ P(x)) (\lambda z. \text{chase}(y, z))$

[Example 10-48](#) may be slightly hard to read at first: you need to see that it



$\lambda z. \text{chase}(y, z)$ . [Example 10-48](#) is equivalent via  $\beta$ -reduction to  $\text{exists } x. (\text{dog}(x) \ \& \ \text{chase}(y, x))$ .

Now let's replace the function expression in [Example 10-48](#) by a variable  $X$  of the same type as an NP, that is, of type  $\langle e, t \rangle$ .

*Example 10-49.*

$X(\lambda z. \text{chase}(y, z))$

The representation of a transitive verb will have to apply to an argument of the type of  $X$  to yield a function expression of the type of VPs, that is, of type  $\langle e, t \rangle$ . We can ensure this by abstracting over both the  $X$  variable in [Example 10-49](#) and also the subject variable  $y$ . So the full solution is reached by giving *chases* the semantic representation shown in [Example 10-50](#).

*Example 10-50.*

$\lambda X \ y. X(\lambda x. \text{chase}(y, x))$

If [Example 10-50](#) is applied to [Example 10-43](#), the result after  $\beta$ -reduction is equivalent to [Example 10-47](#), which is what we wanted all along:

```
>>> lp = nltk.LogicParser()
>>> tvp = lp.parse('X X X (\y.chase(x,y))')
>>> np = lp.parse('P.exists x. (dog(x) & P(x))')
>>> vp = nltk.sem.ApplicationExpression(tvp, np)
>>> print vp
(\X x.X(\y.chase(x,y)) (\P.exists x. (dog(x) & P(x)))
>>> print vp.simplify()
\X.exists z2. (dog(z2) & chase(x,z2))
```

In order to build a semantic representation for a sentence, we also need to combine in the semantics of the subject NP. If the latter is a quantified expression, such as *every girl*, everything proceeds in the same way as we showed for *a dog barks* earlier on; the subject is translated as a function expression which is applied to the semantic representation of the VP. However, we now seem to have created another problem for ourselves with proper names. So far, these have been treated semantically as individual constants, and these cannot be applied as functions to expressions like [Example 10-47](#). Consequently, we need to come up with a different semantic representation for them. What we do in this case is reinterpret proper names so that they too are function expressions, like quantified NPs. Here is the required  $\lambda$ -expression for *Angus*:

*Example 10-51.*

$\lambda P. P(\text{angus})$

[Example 10-51](#) denotes the characteristic function corresponding to the set of all properties which are true of Angus. Converting from an individual constant *angus* to  $\lambda P. P(\text{angus})$  is another example of type-raising, briefly mentioned earlier, and allows us to replace a Boolean-valued application such as  $\lambda x. \text{walk}(x)(\text{angus})$  with an equivalent function application  $\lambda P. P(\text{angus})(\lambda x. \text{walk}(x))$ . By  $\beta$ -reduction, both expressions reduce to  $\text{walk}(\text{angus})$ .

The grammar *simple-sem.fcfg* contains a small set of rules for parsing and translating simple examples of the kind that we have been looking at. Here's a slightly more complicated example:

```
>>> from nltk import load_parser
>>> parser = load_parser('grammars/book_grammars/simple-sem.fcfg')
>>> sentence = 'Angus gives a bone to every dog'
>>> tokens = sentence.split()
>>> trees = parser.nbest_parse(tokens)
>>> for tree in trees:
...     print tree.node['SEM']
all z2. (dog(z2) -> exists z1. (bone(z1) & give(angus,z1,z2)))
```



NLTK provides some utilities to make it easier to derive and inspect semantic interpretations. The function `batch_interpret()` is intended for batch interpretation of a list of input sentences. It builds a dictionary `d` where for each sentence `sent` in the input, `d[sent]` is a list of pairs (*synrep*, *semrep*) consisting of trees and semantic representations for `sent`. The value is a list since `sent` may be syntactically ambiguous; in the following example, however, there is only one parse tree per sentence in the list.

```
(S[SEM=<walk(irene)>]
 (NP[-LOC, NUM='sg', SEM=<\P.P(irene)>]
  (PropN[-LOC, NUM='sg', SEM=<\P.P(irene)>] Irene))
 (VP[NUM='sg', SEM=<\x.walk(x)>]
  (IV[NUM='sg', SEM=<\x.walk(x)>, TNS='pres'] walks)))
(S[SEM=<exists z1.(ankle(z1) & bite(cyril,z1))>]
 (NP[-LOC, NUM='sg', SEM=<\P.P(cyril)>]
  (PropN[-LOC, NUM='sg', SEM=<\P.P(cyril)>] Cyril))
 (VP[NUM='sg', SEM=<\x.exists z1.(ankle(z1) & bite(x,z1))>]
  (TV[NUM='sg', SEM=<\X x.X(\y.bite(x,y))>, TNS='pres'] bites)
  (NP[NUM='sg', SEM=<\Q.exists x.(ankle(x) & Q(x))>]
   (Det[NUM='sg', SEM=<\P Q.exists x.(P(x) & Q(x))>] an)
   (Nom[NUM='sg', SEM=<\x.ankle(x)>]
    (N[NUM='sg', SEM=<\x.ankle(x)>] ankle))))))
```

We have seen now how to convert English sentences into logical forms, and earlier we saw how logical forms could be checked as true or false in a model. Putting these two mappings together, we can check the truth value of English sentences in a given model. Let's take model `m` as defined earlier. The utility `batch_evaluate()` resembles `batch_interpret()`, except that we need to pass a model and a variable assignment as parameters. The output is a triple (*synrep*, *semrep*, *value*), where *synrep*, *semrep* are as before, and *value* is a truth value. For simplicity, the following example only processes a single sentence.

```
>>> v = ""
... bertie => b
... olive => o
... cyril => c
... boy => {b}
... girl => {o}
... dog => {c}
... walk => {o, c}
... see => {(b, o), (c, b), (o, c)}
... ""
>>> val = nltk.parse_valuation(v)
>>> g = nltk.Assignment(val.domain)
>>> m = nltk.Model(val.domain, val)
>>> sent = 'Cyril sees every boy'
>>> grammar_file = 'grammars/book_grammars/simple-sem.fcfg'
>>> results = nltk.batch_evaluate([sent], grammar_file, m, g)[0]
>>> for (syntree, semrep, value) in results:
...     print semrep
...     print value
all z4.(boy(z4) -> see(cyril,z4))
True
```

QUANTIFIER AMBIGUITY REVISITED

One important limitation of the methods described earlier is that they do not deal with scope ambiguity. Our translation method is syntax-driven, in the sense that the semantic representation is closely coupled with the syntactic analysis, and the scope of the quantifiers in the semantics therefore reflects the relative scope of the corresponding NPs in the syntactic parse tree. Consequently, a sentence like

Example 10-26, repeated here, will always be translated as `a` , not `b` .

Example 10-52.

Every girl chases a dog.



2 exists y.(dog(y) & all x.(girl(x) -> chase(x,y)))

There are numerous approaches to dealing with scope ambiguity, and we will look very briefly at one of the simplest. To start with, let’s briefly consider the structure of scoped formulas. [Figure 10-3](#) depicts the way in which the two readings of [Example 10-52](#) differ.




Figure 10-3. Quantifier scopings.

Let’s consider the lefthand structure first. At the top, we have the quantifier corresponding to *every girl*. The  $\phi$  can be thought of as a placeholder for whatever is inside the scope of the quantifier. Moving downward, we see that we can plug in the quantifier corresponding to *a dog* as an instantiation of  $\phi$ . This gives a new placeholder  $\psi$ , representing the scope of *a dog*, and into this we can plug the “core” of the semantics, namely the open sentence corresponding to *x chases y*. The structure on the righthand side is identical, except we have swapped round the order of the two quantifiers.

In the method known as **Cooper storage**, a semantic representation is no longer an expression of first-order logic, but instead a pair consisting of a “core” semantic representation plus a list of **binding operators**. For the moment, think of a binding operator as being identical to the semantic representation of a quantified NP such as [Example 10-44](#) or [Example 10-45](#). Following along the lines indicated in [Figure 10-3](#), let’s assume that we have constructed a Cooper-storage-style semantic representation of sentence [Example 10-52](#), and let’s take our core to be the open formula  $\text{chase}(x,y)$ . Given a list of binding operators corresponding to the two NPs in [Example 10-52](#), we pick a binding operator off the list, and combine it with the core.

```
\P.exists y.(dog(y) & P(y)) (\z2.chase(z1,z2))
```

Then we take the result, and apply the next binding operator from the list to it.

```
\P.all x.(girl(x) -> P(x)) (\z1.exists x.(dog(x) & chase(z1,x)))
```

Once the list is empty, we have a conventional logical form for the sentence. Combining binding operators with the core in this way is called **S-Retrieval**. If we are careful to allow every possible order of binding operators (for example, by taking all permutations of the list; see [Doing More with Functions](#)), then we will be able to generate every possible scope ordering of quantifiers.

The next question to address is how we build up a core+store representation compositionally. As before, each phrasal and lexical rule in the grammar will have a SEM feature, but now there will be embedded features CORE and STORE. To illustrate the machinery, let’s consider a simpler example, namely *Cyril smiles*. Here’s a lexical rule for the verb *smiles* (taken from the grammar *storage.fcfg*), which looks pretty innocuous:



---

```
NP[SEM=[CORE=<@x>, STORE=(<bo(\P.P(cyril),@x)>)] -> 'Cyril'
```

---

The `bo` predicate has two subparts: the standard (type-raised) representation of a proper name, and the expression `@x`, which is called the **address** of the binding operator. (We'll explain the need for the address variable shortly.) `@x` is a metavariable, that is, a variable that ranges over individual variables of the logic and, as you will see, also provides the value of `core`. The rule for `VP` just percolates up the semantics of the `IV`, and the interesting work is done by the `S` rule.

---

```
VP[SEM=?s] -> IV[SEM=?s]

S[SEM=[CORE=<?vp(?np)>, STORE=(?b1+?b2)]] ->
  NP[SEM=[CORE=?np, STORE=?b1]] VP[SEM=[CORE=?vp, STORE=?b2]]
```

---

The `core` value at the `S` node is the result of applying the `VP`'s `core` value, namely `\x.smile(x)`, to the subject `NP`'s value. The latter will not be `@x`, but rather an instantiation of `@x`, say, `z3`. After  $\beta$ -reduction, `<?vp(?np)>` will be unified with `<smile(z3)>`. Now, when `@x` is instantiated as part of the parsing process, it will be instantiated uniformly. In particular, the occurrence of `@x` in the subject `NP`'s `STORE` will also be mapped to `z3`, yielding the element `bo(\P.P(cyril),z3)`. These steps can be seen in the following parse tree.

---

```
(S[SEM=[CORE=<smile(z3)>, STORE=(bo(\P.P(cyril),z3))]])
(NP[SEM=[CORE=<z3>, STORE=(bo(\P.P(cyril),z3))]] Cyril)
(VP[SEM=[CORE=<\x.smile(x)>, STORE=()]]
 (IV[SEM=[CORE=<\x.smile(x)>, STORE=()]] smiles)))
```

---

Let's return to our more complex example, [Example 10-52](#), and see what the storage style SEM value is, after parsing with grammar *storage.fcfg*.

---

```
CORE  = <chase(z1,z2)>
STORE = (bo(\P.all x.(girl(x) -> P(x)),z1), bo(\P.exists x.(dog(x
```

---

It should be clearer now why the address variables are an important part of the binding operator. Recall that during S-retrieval, we will be taking binding operators off the `STORE` list and applying them successively to the `CORE`. Suppose we start with `bo(\P.all x.(girl(x) -> P(x)),z1)`, which we want to combine with `chase(z1,z2)`. The quantifier part of the binding operator is `\P.all x.(girl(x) -> P(x))`, and to combine this with `chase(z1,z2)`, the latter needs to first be turned into a  $\lambda$ -abstract. How do we know which variable to abstract over? This is what the address `z1` tells us, i.e., that *every girl* has the role of chaser rather than chasee.

The module `nltk.sem.cooper_storage` deals with the task of turning storage-style semantic representations into standard logical forms. First, we construct a `CooperStore` instance, and inspect its `STORE` and `CORE`.

---

```
>>> from nltk.sem import cooper_storage as cs
>>> sentence = 'every girl chases a dog'
>>> trees = cs.parse_with_bindops(sentence, grammar='grammars/boo
>>> semrep = trees[0].node['SEM']
>>> cs_semrep = cs.CooperStore(semrep)
>>> print cs_semrep.core
chase(z1,z2)
>>> for bo in cs_semrep.store:
...     print bo
bo(\P.all x.(girl(x) -> P(x)),z1)
bo(\P.exists x.(dog(x) & P(x)),z2)
```

---





```
>>> cs_semrep.s_retrieve(trace=True)
Permutation 1
(\P.all x.(girl(x) -> P(x))) (\z1.chase(z1,z2))
(\P.exists x.(dog(x) & P(x))) (\z2.all x.(girl(x) -> chase(x,z2)
Permutation 2
(\P.exists x.(dog(x) & P(x))) (\z2.chase(z1,z2))
(\P.all x.(girl(x) -> P(x))) (\z1.exists x.(dog(x) & chase(z1,x)

>>> for reading in cs_semrep.readings:
...     print reading
exists x.(dog(x) & all z3.(girl(z3) -> chase(z3,x)))
all x.(girl(x) -> exists z4.(dog(z4) & chase(x,z4)))
```

Discourse Semantics

A **discourse** is a sequence of sentences. Very often, the interpretation of a sentence in a discourse depends on what preceded it. A clear example of this comes from anaphoric pronouns, such as *he*, *she*, and *it*. Given a discourse such as *Angus used to have a dog. But he recently disappeared.*, you will probably interpret *he* as referring to Angus’s dog. However, in *Angus used to have a dog. He took him for walks in New Town.*, you are more likely to interpret *he* as referring to Angus himself.

DISCOURSE REPRESENTATION THEORY

The standard approach to quantification in first-order logic is limited to single sentences. Yet there seem to be examples where the scope of a quantifier can extend over two or more sentences. We saw one earlier, and here’s a second example, together with a translation.

Example 10-54.

- 1. Angus owns a dog. It bit Irene.
- 2.  $\exists x.(dog(x) \ \& \ own(Angus, x) \ \& \ bite(x, Irene))$

That is, the NP *a dog* acts like a quantifier which binds the *it* in the second sentence. Discourse Representation Theory (DRT) was developed with the specific goal of providing a means for handling this and other semantic phenomena which seem to be characteristic of discourse. A **discourse representation structure** (DRS) presents the meaning of discourse in terms of a list of discourse referents and a list of conditions. The **discourse referents** are the things under discussion in the discourse, and they correspond to the individual variables of first-order logic. The **DRS conditions** apply to those discourse referents, and correspond to atomic open formulas of first-order logic.

Figure 10-4 illustrates how a DRS for the first sentence in a is augmented to become a DRS for both sentences.



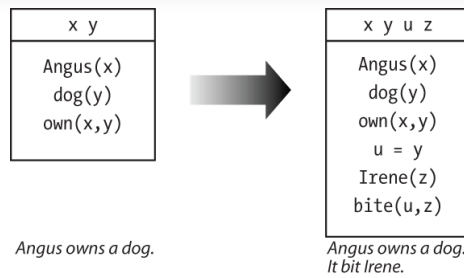


Figure 10-4. Building a DRS: The DRS on the lefthand side represents the result of processing the first sentence in the discourse, while the DRS on the righthand side shows the effect of processing the second sentence and integrating its content.

When the second sentence of `a` is processed, it is interpreted in the context of what is already present in the lefthand side of [Figure 10-4](#). The pronoun *it* triggers the addition of a new discourse referent, say, *u*, and we need to find an **anaphoric antecedent** for it—that is, we want to work out what *it* refers to. In DRT, the task of finding the antecedent for an anaphoric pronoun involves linking it to a discourse referent already within the current DRS, and *y* is the obvious choice. (We will say more about anaphora resolution shortly.) This processing step gives rise to a new condition  $u = y$ . The remaining content contributed by the second sentence is also merged with the content of the first, and this is shown on the righthand side of [Figure 10-4](#).

[Figure 10-4](#) illustrates how a DRS can represent more than just a single sentence. In this case, it is a two-sentence discourse, but in principle a single DRS could correspond to the interpretation of a whole text. We can inquire into the truth conditions of the righthand DRS in [Figure 10-4](#). Informally, it is true in some situation *s* if there are entities **a**, **c**, and **i** in *s* corresponding to the discourse referents in the DRS such that all the conditions are true in *s*; that is, **a** is named *Angus*, **c** is a dog, **a** owns **c**, **i** is named *Irene*, and **c** bit **i**.

In order to process DRSs computationally, we need to convert them into a linear format. Here's an example, where the DRS is a pair consisting of a list of discourse referents and a list of DRS conditions:

```
([x, y], [angus(x), dog(y), own(x,y)])
```

The easiest way to build a DRS object in NLTK is by parsing a string representation [1](#).

```
>>> dp = nltk.DrtParser()
>>> drs1 = dp.parse('([x, y], [angus(x), dog(y), own(x, y)])') 1
>>> print drs1
([x,y],[angus(x), dog(y), own(x,y)])
```

We can use the `draw()` method [1](#) to visualize the result, as shown in [Figure 10-5](#).

```
>>> drs1.draw() 1
```






Figure 10-5. DRS screenshot.

When we discussed the truth conditions of the DRSs in Figure 10-4, we assumed that the topmost discourse referents were interpreted as existential quantifiers, while the conditions were interpreted as though they are conjoined. In fact, every DRS can be translated into a formula of first-order logic, and the `fol()` method implements this translation.

```
>>> print drs1.fol()
exists x y. ((angus(x) & dog(y)) & own(x,y))
```

In addition to the functionality available for first-order logic expressions, DRT Expressions have a DRS-concatenation operator, represented as the `+` symbol. The concatenation of two DRSs is a single DRS containing the merged discourse referents and the conditions from both arguments. DRS-concatenation automatically  $\alpha$ -converts bound variables to avoid name-clashes.

```
>>> drs2 = dp.parse('([x], [walk(x)]) + ([y], [run(y)])')
>>> print drs2
((([x],[walk(x)]) + ([y],[run(y)]))
>>> print drs2.simplify()
([x,y],[walk(x), run(y)])
```

While all the conditions seen so far have been atomic, it is possible to embed one DRS within another, and this is how universal quantification is handled. In `drs3`, there are no top-level discourse referents, and the sole condition is made up of two sub-DRSs, connected by an implication. Again, we can use `fol()` to get a handle on the truth conditions.

```
>>> drs3 = dp.parse('([], [([x], [dog(x)]) -> ([y],[ankle(y), bi
>>> print drs3.fol()
all x.(dog(x) -> exists y.(ankle(y) & bite(x,y)))
```

We pointed out earlier that DRT is designed to allow anaphoric pronouns to be interpreted by linking to existing discourse referents. DRT sets constraints on which discourse referents are “accessible” as possible antecedents, but is not intended to explain how a particular antecedent is chosen from the set of candidates. The module `nltk.sem.drt_resolve_anaphora` adopts a similarly conservative strategy: if the DRS contains a condition of the form `PRO(x)`, the method `resolve_anaphora()` replaces this with a condition of the form `x = [...]`, where `[...]` is a list of possible antecedents.

```
>>> drs4 = dp.parse('([x, y], [angus(x), dog(y), own(x, y)])')
>>> drs5 = dp.parse('([u, z], [PRO(u), irene(z), bite(u, z)])')
>>> drs6 = drs4 + drs5
>>> print drs6.simplify()
([x,y,u,z],[angus(x), dog(y), own(x,y), PRO(u), irene(z), bite(u,
>>> print drs6.simplify().resolve_anaphora()
([x,y,u,z],[angus(x), dog(y), own(x,y), (u = [x,y,z]), irene(z),
```



Our treatment of DRSs is fully compatible with the existing machinery for handling  $\lambda$ -abstraction, and consequently it is straightforward to build compositional semantic representations that are based on DRT rather than first-order logic. This technique is illustrated in the following rule for indefinites (which is part of the grammar *drt.fcfg*). For ease of comparison, we have added the parallel rule for indefinites from *simple-sem.fcfg*.

```
Det[ NUM=sg, SEM=<\P Q. ([ [x], []] + P(x) + Q(x))> ] -> 'a'
Det[ NUM=sg, SEM=<\P Q. exists x. (P(x) & Q(x))> ] -> 'a'
```

To get a better idea of how the DRT rule works, look at this subtree for the NP *a dog*:

```
(NP[ NUM='sg', SEM=<\Q. ([ [x], [dog(x)] ] + Q(x))> ]
  (Det[ NUM='sg', SEM=<\P Q. ([ [x], []] + P(x) + Q(x))> ] a)
  (Nom[ NUM='sg', SEM=<\x. ([ [], [dog(x)] ]> ]
    (N[ NUM='sg', SEM=<\x. ([ [], [dog(x)] ]> dog) ) ) ) )
```

The  $\lambda$ -abstract for the indefinite is applied as a function expression to  $\lambda x. ([[], [dog(x)]]$  which leads to  $\lambda Q. ([ [x], []] + ([[], [dog(x)] ] + Q(x)))$ ; after simplification, we get  $\lambda Q. ([ [x], [dog(x)] ] + Q(x))$  as the representation for the NP as a whole.

In order to parse with grammar *drt.fcfg*, we specify in the call to `load_earley()` that SEM values in feature structures are to be parsed using `DrtParser` in place of the default `LogicParser`.

```
>>> from nltk import load_parser
>>> parser = load_parser('grammars/book_grammars/drt.fcfg', logic
>>> trees = parser.nbest_parse('Angus owns a dog'.split())
>>> print trees[0].node['SEM'].simplify()
([x,z2],[Angus(x), dog(z2), own(x,z2)])
```

DISCOURSE PROCESSING

When we interpret a sentence, we use a rich context for interpretation, determined in part by the preceding context and in part by our background assumptions. DRT provides a theory of how the meaning of a sentence is integrated into a representation of the prior discourse, but two things have been glaringly absent from the processing approach just discussed. First, there has been no attempt to incorporate any kind of inference; and second, we have only processed individual sentences. These omissions are redressed by the module `nltk.inference.discourse`.

Whereas a discourse is a sequence  $s_1, \dots, s_n$  of sentences, a *discourse thread* is a sequence  $s_1-r_1, \dots, s_n-r_n$  of readings, one for each sentence in the discourse. The module processes sentences incrementally, keeping track of all possible threads when there is ambiguity. For simplicity, the following example ignores scope ambiguity:

```
>>> dt = nltk.DiscourseTester(['A student dances', 'Every student
>>> dt.readings()
s0 readings: s0-r0: exists x.(student(x) & dance(x))
s1 readings: s1-r0: all x.(student(x) -> person(x))
```

When a new sentence is added to the current discourse, setting the parameter `consistchk=True` causes consistency to be checked by invoking the model checker for each thread, i.e., each sequence of admissible readings. In this case, the user has the option of retracting the sentence in question.

```
>>> dt.add_sentence('No person dances', consistchk=True)
```



```
>>> dt.retract_sentence('No person dances', verbose=True)
Current sentences are
s0: A student dances
s1: Every student is a person
```

In a similar manner, we use `informchk=True` to check whether a new sentence  $\phi$  is informative relative to the current discourse. The theorem prover treats existing sentences in the thread as assumptions and attempts to prove  $\phi$ ; it is informative if no such proof can be found.

```
>>> dt.add_sentence('A person dances', informchk=True)
Sentence 'A person dances' under reading 'exists x.(person(x) & d
Not informative relative to thread 'd0'
```

It is also possible to pass in an additional set of assumptions as background knowledge and use these to filter out inconsistent readings; see the Discourse HOWTO at <http://www.nltk.org/howto> (<http://www.nltk.org/howto>) for more details.

The `discourse` module can accommodate semantic ambiguity and filter out readings that are not admissible. The following example invokes both Glue Semantics as well as DRT. Since the Glue Semantics module is configured to use the wide-coverage Malt dependency parser, the input (*Every dog chases a boy. He runs.*) needs to be tagged as well as tokenized.

```
>>> from nltk.tag import RegexpTagger
>>> tagger = RegexpTagger(
...     [ ('^(chases|runs)$', 'VB'),
...       ('^(a)$', 'ex_quant'),
...       ('^(every)$', 'univ_quant'),
...       ('^(dog|boy)$', 'NN'),
...       ('^(He)$', 'PRP')
... ])
>>> rc = nltk.DrtGlueReadingCommand(depparser=nltk.MaltParser(tag
>>> dt = nltk.DiscourseTester(['Every dog chases a boy', 'He runs
>>> dt.readings()
s0 readings:
s0-r0: ([], [([x], [dog(x)]) -> ([z3], [boy(z3), chases(x, z3)])])
s0-r1: ([z4], [boy(z4), ([x], [dog(x)]) -> ([], [chases(x, z4)])])

s1 readings:
s1-r0: ([x], [PRO(x), runs(x)])
```

The first sentence of the discourse has two possible readings, depending on the quantifier scoping. The unique reading of the second sentence represents the pronoun *He* via the condition  $PRO(x)$ . Now let's look at the discourse threads that result:

```
>>> dt.readings(show_thread_readings=True)
d0: ['s0-r0', 's1-r0'] : INVALID: AnaphoraResolutionException
d1: ['s0-r1', 's1-r0'] : ([z6, z10], [boy(z6), ([x], [dog(x)]) ->
([], [chases(x, z6)])], (z10 = z6), runs(z10))
```

When we examine threads `d0` and `d1`, we see that reading `s0-r0`, where *every dog* out-scopes *a boy*, is deemed inadmissible because the pronoun in the second sentence cannot be resolved. By contrast, in thread `d1` the pronoun (relettered to `z10`) has been bound via the equation  $(z10 = z6)$ .

Inadmissible readings can be filtered out by passing the parameter `filter=True`.

```
>>> dt.readings(show_thread_readings=True, filter=True)
d1: ['s0-r1', 's1-r0'] : ([z12, z15], [boy(z12), ([x], [dog(x)]) ->
([], [chases(x, z12)])], (z17 = z15), runs(z15))
```



sentences, and also a feel for the techniques that can be deployed to address them.

## Summary

- First-order logic is a suitable language for representing natural language meaning in a computational setting since it is flexible enough to represent many useful aspects of natural meaning, and there are efficient theorem provers for reasoning with first-order logic. (Equally, there are a variety of phenomena in natural language semantics which are believed to require more powerful logical mechanisms.)
- As well as translating natural language sentences into first-order logic, we can state the truth conditions of these sentences by examining models of first-order formulas.
- In order to build meaning representations compositionally, we supplement first-order logic with the  $\lambda$ -calculus.
- $\beta$ -reduction in the  $\lambda$ -calculus corresponds semantically to application of a function to an argument. Syntactically, it involves replacing a variable bound by  $\lambda$  in the function expression with the expression that provides the argument in the function application.
- A key part of constructing a model lies in building a valuation which assigns interpretations to non-logical constants. These are interpreted as either  $n$ -ary predicates or as individual constants.
- An open expression is an expression containing one or more free variables. Open expressions receive an interpretation only when their free variables receive values from a variable assignment.
- Quantifiers are interpreted by constructing, for a formula  $\phi[x]$  open in variable  $x$ , the set of individuals which make  $\phi[x]$  true when an assignment  $g$  assigns them as the value of  $x$ . The quantifier then places constraints on that set.
- A closed expression is one that has no free variables; that is, the variables are all bound. A closed sentence is true or false with respect to all variable assignments.
- If two formulas differ only in the label of the variable bound by binding operator (i.e.,  $\lambda$  or a quantifier), they are said to be  $\alpha$ -equivalents. The result of relabeling a bound variable in a formula is called  $\alpha$ -conversion.
- Given a formula with two nested quantifiers  $Q_1$  and  $Q_2$ , the outermost quantifier  $Q_1$  is said to have wide scope (or scope over  $Q_2$ ). English sentences are frequently ambiguous with respect to the scope of the quantifiers they contain.
- English sentences can be associated with a semantic representation by treating SEM as a feature in a feature-based grammar. The SEM value of a complex expressions, typically involves functional application of the SEM values of the component expressions.

### Further Reading

Consult <http://www.nltk.org/> (<http://www.nltk.org/>) for further materials on this chapter and on how to install the Prover9 theorem prover and Mace4 model builder. General information about these two inference tools is given by (McCune, 2008).

For more examples of semantic analysis with NLTK, please see the semantics and logic HOWTOs at <http://www.nltk.org/howto> (<http://www.nltk.org/howto>). Note that there are implementations of two other approaches to scope ambiguity,



There are many phenomena in natural language semantics that have not been touched on in this chapter, most notably:

- 1. Events, tense, and aspect
- 2. Semantic roles
- 3. Generalized quantifiers, such as *most*
- 4. Intensional constructions involving, for example, verbs such as *may* and *believe*

While (1) and (2) can be dealt with using first-order logic, (3) and (4) require different logics. These issues are covered by many of the references in the following readings.

A comprehensive overview of results and techniques in building natural language front-ends to databases can be found in (Androutsopoulos, Ritchie & Thanisch, 1995).

Any introductory book to modern logic will present propositional and first-order logic. (Hodges, 1977) is highly recommended as an entertaining and insightful text with many illustrations from natural language.

For a wide-ranging, two-volume textbook on logic that also presents contemporary material on the formal semantics of natural language, including Montague Grammar and intensional logic, see (Gamut, 1991a, 1991b). (Kamp & Reyle, 1993) provides the definitive account of Discourse Representation Theory, and covers a large and interesting fragment of natural language, including tense, aspect, and modality. Another comprehensive study of the semantics of many natural language constructions is (Carpenter, 1997).

There are numerous works that introduce logical semantics within the framework of linguistic theory. (Chierchia & McConnell-Ginet, 1990) is relatively agnostic about syntax, while (Heim & Kratzer, 1998) and (Larson & Segal, 1995) are both more explicitly oriented toward integrating truth-conditional semantics into a Chomskyan framework.

(Blackburn & Bos, 2005) is the first textbook devoted to computational semantics, and provides an excellent introduction to the area. It expands on many of the topics covered in this chapter, including underspecification of quantifier scope ambiguity, first-order inference, and discourse processing.

To gain an overview of more advanced contemporary approaches to semantics, including treatments of tense and generalized quantifiers, try consulting (Lappin, 1996) or (van Benthem & ter Meulen, 1997).

Exercises

- 1. ◦ Translate the following sentences into propositional logic and verify that they parse with `LogicParser`. Provide a key that shows how the propositional variables in your translation correspond to expressions of English.
  - 1. If Angus sings, it is not the case that Bertie sulks.
  - 2. Cyril runs and barks.
  - 3. It will snow if it doesn't rain.
  - 4. It's not the case that Irene will be happy if Olive or Tofu comes.
  - 5. Pat didn't cough or sneeze.
  - 6. If you don't come if I call, I won't come if you call.

- 2. ◦ Translate the following sentences into predicate-argument formulas of



2. Tofu is taller than Bertie.
  3. Bruce loves himself and Pat does too.
  4. Cyril saw Bertie, but Angus didn't.
  5. Cyril is a four-legged friend.
  6. Tofu and Olive are near each other.
3. ○ Translate the following sentences into quantified formulas of first-order logic.
1. Angus likes someone and someone likes Julia.
  2. Angus loves a dog who loves him.
  3. Nobody smiles at Pat.
  4. Somebody coughs and sneezes.
  5. Nobody coughed or sneezed.
  6. Bruce loves somebody other than Bruce.
  7. Nobody other than Matthew loves Pat.
  8. Cyril likes everyone except for Irene.
  9. Exactly one person is asleep.
4. ○ Translate the following verb phrases using  $\lambda$ -abstracts and quantified formulas of first-order logic.
1. feed Cyril and give a cappuccino to Angus
  2. be given 'War and Peace' by Pat
  3. be loved by everyone
  4. be loved or detested by everyone
  5. be loved by everyone and detested by no-one

5. ○ Consider the following statements:

```
>>> lp = nltk.LogicParser()
>>> e2 = lp.parse('pat')
>>> e3 = nltk.sem.ApplicationExpression(e1, e2)
>>> print e3.simplify()
exists y.love(pat, y)
```

Clearly something is missing here, namely a declaration of the value of **e1**. In order for `ApplicationExpression(e1, e2)` to be  $\beta$ -convertible to `exists y.love(pat, y)`, **e1** must be a  $\lambda$ -abstract which can take **pat** as an argument. Your task is to construct such an abstract, bind it to **e1**, and satisfy yourself that these statements are all satisfied (up to alphabetic variance). In addition, provide an informal English translation of `e3.simplify()`.

Now carry on doing this same task for the further cases of `e3.simplify()` shown here:

```
>>> print e3.simplify()
exists y.(love(pat,y) | love(y,pat))
```

```
>>> print e3.simplify()
exists y.(love(pat,y) | love(y,pat))
```





```
>>> print e3.simplify()
walk(fido)
```

6. ◦ As in the preceding exercise, find a  $\lambda$ -abstract **e1** that yields results equivalent to those shown here:

```
>>> e2 = lp.parse('chase')
>>> e3 = nltk.sem.ApplicationExpression(e1, e2)
>>> print e3.simplify()
\ x.all y. (dog(y) -> chase(x,pat))
```

```
>>> e2 = lp.parse('chase')
>>> e3 = nltk.sem.ApplicationExpression(e1, e2)
>>> print e3.simplify()
\ x.exists y. (dog(y) & chase(pat,x))
```

```
>>> e2 = lp.parse('give')
>>> e3 = nltk.sem.ApplicationExpression(e1, e2)
>>> print e3.simplify()
\ x0 x1.exists y. (present(y) & give(x1,y,x0))
```

7. ◦ As in the preceding exercise, find a  $\lambda$ -abstract **e1** that yields results equivalent to those shown here:

```
>>> e2 = lp.parse('bark')
>>> e3 = nltk.sem.ApplicationExpression(e1, e2)
>>> print e3.simplify()
exists y. (dog(x) & bark(x))
```

```
>>> e2 = lp.parse('bark')
>>> e3 = nltk.sem.ApplicationExpression(e1, e2)
>>> print e3.simplify()
bark(fido)
```

```
>>> e2 = lp.parse('\ P. all x. (dog(x) -> P(x))')
>>> e3 = nltk.sem.ApplicationExpression(e1, e2)
>>> print e3.simplify()
all x. (dog(x) -> bark(x))
```


8. ◐ Develop a method for translating English sentences into formulas with binary **generalized quantifiers**. In such an approach, given a generalized quantifier  $Q$ , a quantified formula is of the form  $Q(A, B)$ , where both  $A$  and  $B$  are expressions of type  $\langle e, t \rangle$ . Then, for example, **all**( $A, B$ ) is true iff  $A$  denotes a subset of what  $B$  denotes.
9. ◐ Extend the approach in the preceding exercise so that the truth conditions for quantifiers such as *most* and *exactly three* can be computed in a model.
10. ◐ Modify the `sem.evaluate` code so that it will give a helpful error message if an expression is not in the domain of a model's valuation function.
11. • Select three or four contiguous sentences from a book for children. A possible source of examples are the collections of stories in `nltk.corpus.gutenberg:bryant-stories.txt`, `burgess-busterbrown.txt`, and `edgeworth-parents.txt`. Develop a grammar that will allow your sentences to be translated into first-order logic, and build a model that will allow those translations to be checked for truth or falsity.
12. • Carry out the preceding exercise, but use DRT as the meaning representation.



$(P(x) \ \& \ Q(x))$ , convert it to  $(Q(x) \ \& \ P(x))$  if the extension of Q is smaller than the extension of P.

[Support / Sign Out](#)

 [PREV](#)  
[9. Building Feature-Based Grammars](#)

[NEXT](#)   
[11. Managing Linguistic Data](#)

